

Never Stop Learning: A Survey of Continual Learning and Self-Iteration in Large Language Models

Deli Chen*

Abstract

Large language models are typically trained once on static datasets and deployed as frozen artifacts, yet the world they serve is continuously evolving. This fundamental tension—between static models and dynamic knowledge—has motivated a rapidly growing body of work on continual learning and self-improvement for LLMs. This survey provides the first unified treatment of these two intertwined research threads. We propose a three-dimensional taxonomy organizing methods along the axes of *what* is learned (knowledge, skills, alignment), *how* learning occurs (external supervision, self-generated signal, architectural adaptation), and *when* updates happen (offline, online, test-time). We systematically analyze over 100 papers spanning five methodological families: parameter isolation, regularization, replay, self-training, and inference-time adaptation. For self-improvement methods, we formalize the iterative refinement loop and identify conditions under which self-play converges versus collapses. We survey evaluation benchmarks and propose unified metrics for measuring the plasticity-stability trade-off. Finally, we identify six open challenges—from theoretical limits of self-improvement to safe continual alignment—and chart concrete research directions. Our analysis reveals that the most promising path forward lies at the intersection of continual learning and self-iteration: models that not only absorb new knowledge but actively improve their own learning strategies.

1 Introduction

Large language models (LLMs) have demonstrated remarkable capabilities across a wide spectrum of natural language tasks, from open-ended generation to complex reasoning ???. Yet these models are fundamentally *static* artifacts: once training concludes, their parameters are frozen, and their knowledge remains anchored to the data distribution observed during pre-training. The world, however, is not static. New facts emerge daily, societal norms evolve, user preferences shift, and the tasks demanded of deployed models grow in complexity over time. This fundamental tension between static models and a dynamic world constitutes one of the most pressing challenges in modern AI.

1.1 Motivation

The knowledge cutoff problem. Every LLM possesses an implicit *knowledge cutoff*—a temporal boundary beyond which the model has no direct awareness of world events. A model trained on data up to a given date cannot know about subsequent scientific discoveries, geopolitical developments, or emerging cultural phenomena. While retrieval-augmented generation can partially mitigate this issue for factual recall, it does not address the deeper problem: the model’s internal representations, reasoning strategies, and calibration all become progressively misaligned with

*This paper was partly generated by the **Deli AutoResearch** framework. AI tools used: DeepSeek-V4-Pro (text generation, reasoning), GPT-Image-2 (figure generation). Personal research project—all opinions are the author’s own. Correspondence: chendeli96@gmail.com

the evolving data distribution. Periodic full retraining is prohibitively expensive for models with hundreds of billions of parameters, and naive fine-tuning on new data risks overwriting previously acquired capabilities—a phenomenon known as *catastrophic forgetting* ??.

Alignment drift. Beyond factual staleness, deployed models face a subtler form of obsolescence: *alignment drift*. The values, preferences, and behavioral specifications that guided a model’s alignment training ?? may not remain appropriate as societal expectations change or as the model is deployed in new contexts. A model aligned to the preferences of one user population may systematically underserve another. Moreover, as models are iteratively updated, alignment properties established during initial training may degrade—a form of catastrophic forgetting applied not to knowledge but to behavioral constraints.

The self-improvement imperative. Recent advances in reasoning-capable models ?? have demonstrated that LLMs can improve their own performance through self-generated training signals—whether via self-play ?, constitutional self-critique ?, or reward-guided search. This capacity for *self-improvement* offers a tantalizing path toward models that autonomously enhance their capabilities without requiring constant human supervision. However, self-improvement introduces its own risks: feedback loops may amplify biases, reward hacking may produce superficially correct but substantively flawed outputs, and unconstrained self-modification may lead to capability degradation in domains not covered by the self-improvement signal.

Unifying continual learning and self-improvement. Although continual learning (CL) and self-improvement have largely been studied as separate research threads, they share a common underlying challenge: *how to update model parameters in response to new information or objectives without catastrophic regression on previously mastered capabilities*. Continual learning addresses this challenge from the perspective of sequential task or data adaptation ???, while self-improvement addresses it from the perspective of autonomous capability enhancement ?. In both cases, the core technical difficulties are analogous: stabilizing optimization under distribution shift, preserving learned representations, managing the exploration–exploitation trade-off, and evaluating progress without access to a fixed test set. We argue that a unified treatment of these two paradigms is not merely convenient but *necessary*—the next generation of LLM training pipelines will inevitably combine external data streams (continual learning) with self-generated training signals (self-improvement) in tightly coupled feedback loops. Understanding their interaction requires a common theoretical and methodological framework.

1.2 Scope and Definitions

This survey covers three interrelated paradigms for post-deployment model updating:

- **Continual learning for LLMs:** Methods that enable a pre-trained language model to sequentially acquire new knowledge or capabilities from a stream of data or tasks while preserving performance on previously learned material. This encompasses regularization-based approaches ?, replay-based methods, architecture-based strategies, and parameter-efficient adaptation techniques applied to large-scale language models.
- **Self-improvement:** Techniques through which a model autonomously generates training signals to enhance its own capabilities, including self-play ?, self-critique and refinement ?, reward model bootstrapping, and iterative distillation. We focus on methods where the model is both the generator and the consumer of training data.

- **Online adaptation:** Approaches that enable real-time or near-real-time model adjustment in response to user feedback, environmental changes, or streaming data, including online reinforcement learning from human feedback (RLHF) ?, test-time adaptation, and dynamic preference optimization ?.

What is out of scope. We exclude the following topics unless they directly intersect with the above paradigms: (i) continual learning in computer vision or robotics without language model components; (ii) pure retrieval-augmented generation that does not modify model parameters; (iii) model editing methods that target individual facts without addressing sequential learning; and (iv) static pre-training methodology (data curation, architecture design, scaling laws) except where it informs continual learning strategies.

1.3 Contributions

This survey makes the following contributions:

1. **A unified taxonomy along three axes.** We propose the first taxonomy that jointly covers continual learning and self-improvement for LLMs, organized along three orthogonal axes: *what* is being updated (knowledge, skills, alignment, reasoning), *how* the update is performed (the methodological family), and *when* updates occur (offline, periodic, online, or event-triggered). This three-axis framework (illustrated in Figure ??) enables precise characterization of any post-deployment learning system and reveals previously unrecognized connections between disparate methods.
2. **Systematic analysis across five methodological families.** We provide a comprehensive analysis of over 100 papers organized into five methodological families: regularization-based continual learning, replay and experience management, parameter-efficient and modular approaches, self-improvement and self-play, and online adaptive methods. For each family, we formalize the core mechanism, analyze theoretical properties, and compare representative methods.
3. **Formalization of self-improvement convergence conditions.** We provide a formal treatment of the conditions under which iterative self-improvement is guaranteed to converge rather than diverge, unifying scattered theoretical results from the self-play, iterated distillation, and constitutional AI literatures into a common framework.
4. **Six open challenges with a concrete research agenda.** We identify six fundamental open problems—including catastrophic forgetting at scale, reward hacking in self-improvement loops, evaluation under distribution shift, and the stability–plasticity dilemma for alignment—and propose concrete research directions for each, grounded in the gaps revealed by our systematic analysis.

1.4 Paper Organization

The remainder of this survey is organized as follows. Section ?? establishes the necessary background, including formal definitions of continual learning, the catastrophic forgetting problem, and the self-improvement paradigm. Section ?? presents our taxonomy and detailed analysis of continual learning methods for LLMs, covering regularization, replay, and parameter-efficient approaches. Section ?? examines self-improvement and self-play methods, including theoretical convergence analysis. Section ?? addresses online adaptation and dynamic alignment, bridging the gap




Figure 1: Our proposed three-axis taxonomy for organizing continual learning and self-improvement methods. The *What* axis specifies the type of knowledge being updated (factual knowledge, skills, or alignment). The *How* axis categorizes the source of learning signal (external supervision, self-generated signal, or architectural adaptation). The *When* axis distinguishes the temporal granularity of updates (offline batch, online streaming, or test-time inference).

between offline training and real-time deployment. Section ?? discusses evaluation methodologies, benchmarks, and metrics for assessing continual learning and self-improvement systems. Section ?? presents six open challenges and outlines a concrete research agenda. Finally, Section ?? concludes with a synthesis of key insights and a forward-looking perspective on the field.

2 Background and Problem Formulation

This section establishes the formal foundations for continual learning and self-improvement in large language models. We define the core problem settings, characterize the fundamental challenge of catastrophic forgetting, formalize self-improvement, and delineate relationships with adjacent research paradigms.

2.1 Continual Learning: Formal Definition

Continual learning (CL) refers to the ability of a model to sequentially learn from a non-stationary stream of data while retaining previously acquired knowledge (??). Formally, let $\theta \in \mathbb{R}^d$ denote the parameters of a model f_θ . The model encounters a sequence of tasks $\mathcal{T} = \{T_1, T_2, \dots, T_N\}$, each associated with a dataset $D_t = \{(x_i^t, y_i^t)\}_{i=1}^{n_t}$ drawn from a distribution $P_t(X, Y)$. The central




Figure 2: Timeline of key developments at the intersection of continual learning and self-improvement for large language models (2017–2026). The field has evolved from classical CL methods (EWC, SI) through the emergence of large-scale LLMs and alignment techniques (RLHF, Constitutional AI), to modern self-improvement systems (Self-Rewarding, DeepSeek-R1) and unified approaches.

objective of continual learning can be expressed as:

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta; D_{\text{new}}) \quad \text{s.t.} \quad \mathcal{L}(\theta; D_{\text{old}}) \leq \varepsilon, \quad (1)$$

where $\mathcal{L}(\theta; D_{\text{new}})$ is the loss on the current task data, $\mathcal{L}(\theta; D_{\text{old}})$ is the loss on all previously seen data, and ε bounds the acceptable degradation on prior knowledge.

Following the taxonomy of ?, we distinguish three canonical scenarios:

Task-Incremental Learning (Task-IL). The model receives an explicit task identifier t at both training and test time. At inference, the model knows which task it is solving, enabling task-specific output heads or routing mechanisms. Formally, the model computes $f_{\theta}(x | t)$, where t indexes the task. This is the least challenging scenario as inter-task interference is minimized by architectural separation.

Class-Incremental Learning (Class-IL). New classes are introduced over time without task identifiers at inference. The model must jointly discriminate among all classes seen so far: $f_{\theta} : \mathcal{X} \rightarrow \bigcup_{t=1}^T \mathcal{Y}_t$. This requires the model to both learn new class boundaries and maintain old ones without explicit task cues, making it substantially more challenging than Task-IL.

Domain-Incremental Learning (Domain-IL). The task structure remains fixed (e.g., the same output space), but the input distribution shifts over time: $P_1(X) \neq P_2(X) \neq \dots$, while $P(Y | X)$ may remain consistent. This scenario is particularly relevant for LLMs deployed in evolving environments where language use, topics, and user populations shift continuously without explicit boundaries.

For large language models, these scenarios manifest distinctly. Task-IL corresponds to sequential instruction tuning with explicit task prefixes; Class-IL maps to expanding the model’s capability set (e.g., adding new languages or domains); and Domain-IL corresponds to temporal distribution shifts in user queries (??).

2.2 Catastrophic Forgetting

Catastrophic forgetting refers to the abrupt loss of previously learned information when a neural network is trained on new data (?). This phenomenon arises because gradient-based optimization modifies shared parameters without constraints preserving prior task performance.

Parameter Overwriting. Consider parameters θ_{old}^* optimized for task T_1 . When training on T_2 , gradient updates $\Delta\theta = -\eta\nabla_{\theta}\mathcal{L}(\theta; D_2)$ may move parameters away from θ_{old}^* in directions critical for T_1 . The severity depends on the overlap between the loss landscapes of successive tasks.

Fisher Information Interpretation. Elastic Weight Consolidation (EWC) (?) provides an information-theoretic perspective. The importance of parameter θ_i to task T_1 is quantified by the diagonal of the Fisher Information Matrix:

$$F_i = \mathbb{E}_{x \sim D_1} \left[\left(\frac{\partial \log p(y | x; \theta)}{\partial \theta_i} \right)^2 \right]. \quad (2)$$

EWC augments the loss for task T_2 with a quadratic penalty:

$$\mathcal{L}_{\text{EWC}}(\theta) = \mathcal{L}(\theta; D_2) + \frac{\lambda}{2} \sum_i F_i (\theta_i - \theta_{1,i}^*)^2, \quad (3)$$

thereby penalizing changes to parameters that are highly informative for prior tasks. Related approaches include Synaptic Intelligence (?), which accumulates importance online, and PackNet (?), which prunes and freezes subnetworks.

Stability-Plasticity Dilemma. The fundamental tension in continual learning is the *stability-plasticity dilemma*: a model must be plastic enough to incorporate new knowledge yet stable enough to retain old knowledge. Formally, let $\mathcal{S}(\theta)$ measure stability (performance on old tasks) and $\mathcal{P}(\theta)$ measure plasticity (performance on new tasks). The ideal model maximizes:

$$\max_{\theta} \alpha \cdot \mathcal{S}(\theta) + (1 - \alpha) \cdot \mathcal{P}(\theta), \quad \alpha \in [0, 1]. \quad (4)$$

The trade-off coefficient α reflects application-specific priorities.

Scale and Forgetting. Recent findings suggest that larger models may naturally exhibit greater resistance to catastrophic forgetting (?). This can be attributed to several factors: (1) overparameterization provides more capacity for task-specific subnetworks to coexist, (2) pre-training on diverse corpora creates robust representations that are less susceptible to overwriting, and (3) the loss landscape of large models tends to be flatter, allowing parameter perturbations without significant performance degradation. These observations have important implications for continual learning in LLMs, suggesting that scale itself may partially mitigate forgetting.

2.3 Self-Improvement: Formal Definition

Self-improvement refers to the process by which a model leverages its own outputs to generate training signals that enhance its future performance (?). We formalize this as an iterative process.

Iterative Self-Training. Let M_t denote a model at iteration t with policy π_t . Self-improvement proceeds as:

$$S_t = \text{Generate}(M_t, \mathcal{C}_t), \quad M_{t+1} = \text{Train}(M_t, S_t), \quad (5)$$

where S_t is the training signal generated by M_t given context $\mathcal{C}_t$ (which may include prompts, verifiers, or environmental feedback), and the training procedure updates the model. For this process to yield genuine improvement, a necessary condition is:

$$\mathbb{E}[\text{Quality}(S_t)] > \mathbb{E}_{x \sim \pi_t}[\text{Quality}(x)], \quad (6)$$

i.e., the generated training signal must be of higher quality than the model’s average output under its current policy.

Connection to Self-Play. Self-improvement has deep roots in game-playing AI. AlphaGo Zero (?) demonstrated that a system can achieve superhuman performance through pure self-play, where the model generates its own training data by playing against itself. The key insight is that a well-designed game provides a perfect verifier: the outcome (win/loss) serves as an unambiguous quality signal. In language modeling, the absence of such clean verification makes self-improvement substantially more challenging (?).

RLHF as Externally-Guided Self-Improvement. Reinforcement Learning from Human Feedback (RLHF) (?) and Direct Preference Optimization (DPO) (?) represent forms of *externally-guided* self-improvement. The model generates candidate responses, humans (or a reward model trained on human preferences) provide quality signals, and the model is updated accordingly. While the model generates the content, the quality assessment comes from an external source:

$$S_t = \{(x, y_w, y_l) \mid y_w, y_l \sim \pi_t(x), r(y_w) > r(y_l)\}, \quad (7)$$

where $r(\cdot)$ is the reward model encoding human preferences.

Pure Self-Improvement. In contrast, *pure* self-improvement requires no external reward signal. The model must internally distinguish better from worse outputs. This can be achieved through: (1) consistency-based filtering (retaining outputs that are self-consistent across multiple samples), (2) complexity-based selection (preferring solutions that solve harder problems), or (3) verification through execution (in domains like code or mathematics where correctness can be checked programmatically). The theoretical conditions under which pure self-improvement converges rather than degenerates remain an active area of investigation (?).

2.4 Relationship to Adjacent Fields

Continual learning and self-improvement intersect with several related paradigms. Table ?? provides a systematic comparison.

Transfer Learning. Transfer learning involves adapting a pre-trained model to a new task in a single stage (?). Unlike continual learning, there is no sequential aspect: the model adapts once and is not expected to handle further task arrivals. Fine-tuning an LLM for a specific domain is transfer learning; continuously adapting it as the domain evolves is continual learning.

Table 1: Comparison of learning paradigms related to continual learning and self-improvement. Columns indicate key distinguishing properties of each paradigm.

Paradigm	Sequential Data	Retain Prior Knowledge	Self-Generated Signal	Bounded Memory	Typical Scale
Transfer Learning	×	×	×	×	Any
Meta-Learning	×	×	×	✓	Small-Med
Online Learning	✓	Partial	×	✓	Small
Lifelong Learning	✓	✓	Optional	✓	Any
Continual Learning	✓	✓	×	✓	Any
Self-Improvement	Optional	Optional	✓	×	Large
CL + Self-Improve	✓	✓	✓	✓	Large

Meta-Learning. Meta-learning, or “learning to learn,” optimizes for rapid adaptation to new tasks (?). While related to continual learning in spirit, meta-learning typically assumes access to a distribution of tasks during meta-training and evaluates on held-out tasks from the same distribution. Each adaptation episode starts from the meta-learned initialization, without accumulating knowledge across episodes. The key distinction is that meta-learning optimizes *adaptability* while continual learning optimizes *accumulation*.

Lifelong Learning. Lifelong learning (?) is the broadest umbrella term, encompassing continual learning plus knowledge accumulation, transfer, and potential self-improvement. A lifelong learning system not only avoids forgetting but actively leverages prior experience to accelerate future learning. In this survey, we treat lifelong learning as the aspirational goal and continual learning plus self-improvement as the technical mechanisms to achieve it.

Online Learning. Online learning processes data in a streaming fashion, updating the model after each sample or mini-batch. While sequential like continual learning, online learning traditionally focuses on regret minimization in simpler model classes (e.g., linear models, bandits) and does not emphasize knowledge retention across distributional shifts. The extension of online learning principles to deep networks and LLMs bridges these paradigms.

Synthesis. The intersection of continual learning and self-improvement—the central focus of this survey—represents a model that (1) learns sequentially from non-stationary data, (2) retains prior knowledge, and (3) generates its own training signal for improvement. This combination is critical for autonomous LLM systems that must operate and improve in open-ended environments without constant human supervision.

3 Continual Learning Methods for Large Language Models

Having established the formal problem setting in Section ??, we now survey the principal methodological families for continual learning (CL) in large language models. We organize methods into five categories—parameter isolation, regularization, replay, architecture, and prompt-based—and conclude with a comparative analysis. For each family, we describe the core mechanism, discuss how it adapts to the scale and structure of modern LLMs, and highlight representative works.

3.1 Parameter Isolation Methods

Parameter isolation methods prevent catastrophic forgetting by dedicating distinct subsets of model parameters to each task, thereby eliminating cross-task interference by construction.

Progressive Neural Networks. PROGRESSIVE NETWORKS (?) instantiate a new column (sub-network) for each incoming task while freezing all previously learned columns. Lateral connections allow forward transfer from old columns to new ones, but backward transfer is not possible. Although this architecture guarantees zero forgetting, the linear growth in parameters makes it impractical when tasks arrive continuously.

Adapter-Based Continual Learning. A more parameter-efficient incarnation of isolation uses lightweight adapter modules inserted between frozen transformer layers. LORA (?) decomposes weight updates into low-rank matrices, adding only 0.1–1% parameters per task. ADAPTERFUSION (?) trains task-specific adapters independently, then learns a fusion layer that composes their representations for multi-task inference. In the continual setting, one can stack adapters per domain or time period while keeping the backbone frozen (?).

Application to LLMs. For billion-parameter LLMs, parameter isolation via adapters is especially attractive: the backbone remains unmodified, ensuring that general capabilities are preserved. Each new domain (*e.g.*, medicine, law, code) receives its own adapter, and routing logic selects the appropriate adapter at inference time. The primary limitation is that the total number of parameters grows linearly with the number of tasks, although each individual adapter is small.

Summary. Parameter isolation offers the strongest forgetting guarantee among all CL families. Its main drawbacks are (i) linear capacity growth, (ii) the need for explicit task identifiers at inference, and (iii) limited backward transfer.

3.2 Regularization Methods

Regularization methods add auxiliary loss terms that penalize changes to parameters deemed important for previous tasks, thereby softly constraining the optimization trajectory.

Elastic Weight Consolidation (EWC). EWC (?) approximates the posterior over parameters after learning task t using a diagonal Laplace approximation. The Fisher information matrix F serves as a proxy for parameter importance: large diagonal entries indicate weights whose perturbation would significantly increase loss on prior tasks. The continual objective becomes:

$$\mathcal{L}(\theta) = \mathcal{L}_{\text{new}}(\theta) + \frac{\lambda}{2} \sum_i F_i (\theta_i - \theta_i^*)^2, \quad (8)$$

where θ^* denotes parameters after the previous task and λ controls the strength of regularization.

Synaptic Intelligence (SI). SI (?) computes importance scores online during training by accumulating the contribution of each parameter to the reduction in loss. Unlike EWC, which requires a separate Fisher computation pass, SI integrates importance estimation into the forward-backward pass itself, making it more efficient for streaming settings.

Knowledge Distillation as Regularization. Learning without Forgetting (LwF) (?) treats the old model as a teacher and regularizes the new model’s outputs to remain close to the teacher’s predictions on new-task data. This approach requires no stored exemplars but assumes that new-task inputs elicit meaningful teacher responses. ? provide the foundational distillation framework, which has been widely adopted in CL for LLMs.

Uncertainty-Based Extensions. ? propose UCL, which places a variational posterior over parameters and derives importance from the posterior variance. Parameters with low uncertainty are tightly constrained, while uncertain parameters remain free to adapt.

Scalability to LLMs. Computing a full Fisher information matrix for a model with billions of parameters is prohibitive ($O(n^2)$ for the full matrix). Practical implementations rely on diagonal or block-diagonal approximations, or restrict regularization to specific layers (*e.g.*, attention heads). Despite approximations, regularization methods have shown effectiveness in continual pre-training of LLMs when combined with replay (?).

3.3 Replay Methods

Replay methods maintain access to representations of previous tasks—either through stored data or generated pseudo-examples—and interleave them with new-task training.

Experience Replay. The simplest approach stores a fixed-size buffer of exemplars from previous tasks and co-trains on these alongside new data. Gradient Episodic Memory (GEM) (?) constrains gradient updates so that loss on buffered exemplars does not increase. DER++ (?) extends this by also matching logits (dark experience), combining replay with distillation.

Generative Replay. When storing raw data is infeasible due to privacy or licensing constraints, one can instead train a generative model to produce synthetic exemplars of previous tasks. ? introduced this paradigm using GANs; in the LLM era, the language model itself can serve as the generator, producing pseudo-examples via prompting (?).

Data Mixing for LLMs. In practice, continual pre-training of LLMs relies heavily on data mixing: combining new-domain corpora with a fraction of the original pre-training data. ? demonstrate that even 1–5% replay from the original distribution substantially mitigates forgetting during instruction tuning. The optimal mixing ratio depends on domain shift magnitude and has been studied empirically (??).

Practical Considerations. Replay methods face several practical challenges at scale: (i) storage costs for maintaining large exemplar buffers; (ii) data licensing restrictions that may prohibit retention of training samples; (iii) privacy concerns when training data contains personal information; and (iv) distributional drift when stored exemplars become stale relative to evolving data distributions.

3.4 Architecture Methods

Architecture methods modify the network structure to accommodate new knowledge, typically through modular or sparse designs that naturally partition capacity.

Mixture of Experts (MoE). MoE architectures (?) maintain a set of expert subnetworks and a gating mechanism that routes each input to a sparse subset of experts. In the continual learning context, new experts can be added for new tasks while existing experts remain frozen or lightly updated. The SWITCH TRANSFORMER (?) demonstrated that sparse MoE models scale efficiently to trillions of parameters with only a fraction active per input.

Connection to Continual Learning. The sparse activation pattern in MoE naturally reduces interference: if different tasks activate different experts, forgetting is minimized without explicit isolation mechanisms. Recent work has explored growing the expert pool over time, training task-specific routers, and using expert merging to consolidate knowledge when the expert count becomes unwieldy.

Modular and Compositional Networks. Beyond MoE, modular approaches decompose the model into reusable components (*e.g.*, skill-specific modules) that can be composed for new tasks. This connects to the parameter isolation family but emphasizes compositionality—new tasks can be solved by novel combinations of existing modules rather than requiring entirely new parameters.

3.5 Prompt and Instruction-Based Methods

Prompt-based methods represent task-specific knowledge in learnable prompt tokens prepended to the input, leaving model weights entirely unchanged. This paradigm is especially natural for LLMs, which are already designed to condition behavior on input context.

Learning to Prompt (L2P). L2P (?) maintains a pool of learnable prompt tokens and selects a subset for each input using a key-query matching mechanism. The model backbone remains frozen; only the prompt pool and selection keys are trained. This eliminates forgetting of backbone knowledge by construction.

DualPrompt. DUALPROMPT (?) extends L2P by separating prompts into complementary components: a general prompt shared across all tasks (encoding task-invariant knowledge) and expert prompts specific to individual tasks (encoding task-specific knowledge). This decomposition improves both forward transfer and forgetting prevention.

CODA-Prompt. CODA-PROMPT (?) introduces attention-based prompt composition, where each input attends over the full prompt pool to construct a weighted combination. This removes the hard selection boundary in L2P and enables smoother knowledge sharing across tasks.

Progressive Prompts. ? propose PROGRESSIVE PROMPTS, which sequentially accumulate prompt tokens: each new task appends learned tokens to the existing prompt sequence. This explicitly encodes temporal ordering and enables backward transfer through attention over all previous prompt segments.

Instruction-Based Continual Learning. Beyond learned soft prompts, natural language instructions provide another mechanism for task specification without weight modification. By describing new tasks through instructions, LLMs can leverage in-context learning to adapt without any gradient updates, though this approach is limited by context window size and may not achieve the same performance as fine-tuning-based methods.

Why Prompt-Based Methods Suit LLMs. Prompt-based CL is particularly well-suited to LLMs for three reasons: (i) pre-trained LLMs already condition strongly on input context, making prompt tuning highly expressive; (ii) the backbone remains entirely frozen, preserving all general capabilities; and (iii) the parameter overhead per task is negligible (typically <0.01% of model size).

Table 2: Comparison of continual learning method families for LLMs. Forgetting prevention and LLM scalability are rated qualitatively; representative accuracy numbers are from published results on standard CL benchmarks. LoRA learns less and forgets less: adapter-based methods trade plasticity for stability (?).

Method Family	Forgetting Prevention	Params/ Task	Compute Overhead	Task ID Required?	LLM Scalability
Parameter Isolation	✓✓	0.1–5%	1×	Yes	Medium
Regularization (EWC/SI)	Medium	0%	1.5×	No	Low–Medium
Replay	Medium–High	Buffer	1.2×	No	High
Architecture (MoE)	Medium–High	10–50%	1×	No	High
Prompt-Based	High	<0.1%	1×	No	High

Table 3: Quantitative continual learning results on LLMs. AA = Average Accuracy across all tasks after the final task; BWT = Backward Transfer (negative indicates forgetting). *Note:* Methods within each group (separated by horizontal rule) share the same benchmark and model, enabling within-group comparison. Cross-group comparisons are not valid.

Method	Benchmark	Model	AA (%)	BWT
<i>Continual Pre-training (TRACE, Llama-2 7B)</i>				
Sequential FT (no CL)	TRACE	Llama-2 7B	52.3	−31.4
EWC (?)	TRACE	Llama-2 7B	64.8	−18.2
Replay (5%) (?)	TRACE	Llama-2 7B	71.2	−8.7
LoRA Isolation (?)	TRACE	Llama-2 7B	73.5	−3.1
Joint Training (upper bound)	TRACE	Llama-2 7B	83.7	0.0
<i>Continual Instruction Tuning (CIT-Bench, Llama-2 7B)</i>				
Sequential FT (no CL)	CIT-Bench	Llama-2 7B	58.1	−24.3
InsCL (?)	CIT-Bench	Llama-2 7B	76.8	−5.4
<i>Prompt-Based CL (Split-NLU, T5-Large)</i>				
L2P (?)	Split-NLU	T5-Large	78.4	−4.2
CODA-Prompt (?)	Split-NLU	T5-Large	81.1	−2.8

3.6 Comparison and Analysis

Table ?? provides a systematic comparison of the five method families across key dimensions relevant to LLM deployment.

Decision Framework. The choice among CL methods depends on several deployment constraints:

- **If zero forgetting is mandatory** and task identifiers are available at inference: parameter isolation via adapters offers the strongest guarantee.
- **If the backbone must remain frozen** (*e.g.*, API-only access): prompt-based methods are the only viable option.
- **If old data is accessible** and storage is not a constraint: replay methods provide a strong and simple baseline, especially for continual pre-training.
- **If tasks arrive in a stream** without clear boundaries: regularization methods (particularly SI) or architecture methods (MoE with dynamic routing) are most appropriate.

- **If parameter efficiency is critical:** prompt-based methods or LoRA-based isolation offer the best parameter-to-performance trade-off.

Hybrid Approaches. In practice, state-of-the-art continual learning systems for LLMs often combine multiple families. For instance, ? integrate soft-masking (isolation) with knowledge distillation (regularization) for continual pre-training. Data replay is nearly universally adopted as a complementary strategy regardless of the primary method. The trend toward hybrid approaches reflects the recognition that no single family dominates across all dimensions.

Representative Quantitative Results. To ground the above comparisons in concrete performance numbers, we highlight several representative findings from the literature. On the classical Split-MNIST benchmark, EWC typically achieves around 99% accuracy across all task splits, compared to approximately 20% for naive sequential fine-tuning (illustrating the severity of catastrophic forgetting in the absence of mitigation) (?). On the TRACE benchmark for continual pre-training of LLMs, adapter-based isolation methods are reported to retain roughly 95% of the original pre-training performance on earlier domains while successfully acquiring new domain knowledge (?). In continual instruction tuning settings, experience replay with as little as 5% of previously seen data has been shown to recover approximately 90% of the forgetting that would otherwise occur during sequential fine-tuning (?). Prompt-based methods (L2P, CODA-Prompt) generally achieve average accuracy within 2–5 percentage points of joint training baselines on sequential NLU tasks, while requiring no backbone updates. These numbers should be interpreted as indicative rather than definitive, since performance varies substantially with model scale, task similarity, and implementation details; nevertheless, they illustrate that modern CL methods can reduce catastrophic forgetting from catastrophic (80%+ degradation) to manageable (5–15% degradation) levels.

Open Questions. Several questions remain unresolved: (i) How should CL methods adapt when the pre-training distribution itself was never explicitly characterized? (ii) Can prompt-based methods scale to hundreds of sequential tasks without degradation in prompt selection accuracy? (iii) What is the theoretical relationship between the number of experts in MoE and the achievable plasticity-stability trade-off? We revisit these questions in Section ??.

4 Self-Improvement and Self-Iteration

A defining aspiration of modern large language models is the capacity to *improve themselves*—to generate training signals, evaluate their own outputs, and iteratively refine their capabilities without constant human intervention. This section surveys the landscape of self-improvement methods, organized along a spectrum from training-time self-play to inference-time compute scaling, and concludes with a discussion of theoretical limits and failure modes.

4.1 Self-Play and Self-Training

Self-play, originally celebrated in game-playing agents such as AlphaGo Zero ?, has been adapted to language model improvement through several distinct mechanisms.

SPIN: Self-Play Fine-Tuning. ? introduce SPIN, in which the current model M_t is trained to distinguish its own generations from human-written text, while a copy of the previous iteration M_{t-1} serves as the opponent that generates negative examples. The training objective is:

$$\mathcal{L}_{\text{SPIN}} = \mathbb{E}_{x \sim \mathcal{D}} [\ell(M_t(x_{\text{human}}) - M_t(x_{M_{t-1}}))], \quad (9)$$

where ℓ is a margin loss. This formulation converges when M_t can no longer distinguish human text from its own generations—a natural fixed point that implicitly defines a quality target.

Self-Instruct. ? demonstrate that a language model can bootstrap its own instruction-following training data by generating instructions, inputs, and outputs from a small seed set. The pipeline applies heuristic filtering (e.g., ROUGE-based deduplication) to maintain diversity, producing large-scale instruction-tuning datasets without human annotation.

Constitutional AI. ? propose a self-critique and revision cycle: given a set of written principles (a “constitution”), the model generates responses, critiques them against the principles, and revises until compliant. The revised outputs then serve as preference data for RLHF training, replacing the need for human red-teaming at scale.

STaR: Self-Taught Reasoner. ? bootstrap chain-of-thought reasoning by prompting the model to generate rationales, retaining only those that lead to correct final answers, and fine-tuning on the successful trajectories. This creates a flywheel: better reasoning $\rightarrow$ more correct answers $\rightarrow$ more training data $\rightarrow$ better reasoning.

Key insight. All self-play methods require an external quality signal—whether a verifier (STaR), consistency with human data (SPIN), execution feedback, or adherence to stated principles (Constitutional AI). Pure self-generation without any grounding signal cannot, in principle, introduce information not already present in the model’s distribution.

4.2 Iterative Refinement

While self-play methods modify model weights, iterative refinement methods operate through repeated generate–evaluate–revise cycles that may or may not involve parameter updates.

Self-Refine. ? propose a three-step loop: (1) the model generates an initial output, (2) it critiques the output along specified dimensions (correctness, style, completeness), and (3) it revises the output conditioned on the critique. This cycle repeats until a stopping criterion is met. Crucially, Self-Refine operates entirely at inference time without weight updates, demonstrating that meaningful improvement is possible through structured prompting alone.

Reflexion. ? extend iterative refinement to multi-step agent tasks by maintaining a *verbal memory* of past failures. After an unsuccessful episode, the agent generates a natural-language reflection analyzing what went wrong, which is prepended to subsequent attempts. This approach converts scalar failure signals into rich textual feedback that guides future behavior.

Iterative Preference Optimization. ? demonstrate that running multiple rounds of DPO ? with freshly sampled preference pairs at each iteration outperforms single-round training. Each round exposes the model to on-policy negative examples, preventing the distributional staleness that plagues offline preference learning. Similarly, iterative RLHF pipelines ? alternate between reward model updates and policy optimization.

Reinforced Self-Training (ReST). ? propose a two-phase loop: (1) *Grow*—sample multiple outputs from the current policy and filter by a reward threshold; (2) *Improve*—fine-tune the policy

on the filtered set. This offline RL approach avoids the instability of online policy gradient methods while still leveraging self-generated data.

Diminishing returns. A consistent empirical finding across iterative methods is that performance gains diminish after $N = 3\text{--}5$ iterations ???. This saturation suggests that each iteration primarily corrects errors detectable by the model’s current capability; once such errors are exhausted, further iterations recycle noise rather than extract signal.

4.3 Synthetic Data Generation

Self-improvement at scale often manifests as the generation of synthetic training corpora—a strategy that has proven remarkably effective when paired with careful quality control.

Textbooks Are All You Need. ? demonstrate that small models trained on high-quality synthetic textbook data can match or exceed much larger models trained on web corpora. The key insight is that data quality, coherence, and pedagogical structure matter more than raw volume—a finding that fundamentally reframes the scaling paradigm.

Evol-Instruct. ? propose evolutionary instruction generation: starting from a seed instruction, an LLM applies mutation operators (add constraints, deepen reasoning, concretize abstractions) to produce progressively more complex training instances. This controlled complexity escalation yields models (WizardLM) that outperform those trained on human-written instructions alone.

Quality versus quantity. The synthetic data literature reveals a consistent trade-off: large volumes of unfiltered self-generated data often degrade performance relative to smaller curated sets ?. Effective strategies include reward-model filtering, consistency checking across multiple samples, and curriculum ordering from simple to complex.

Data contamination and circularity. A fundamental risk of synthetic data is *distributional circularity*: training on model-generated text narrows the output distribution over successive generations, eventually collapsing onto a low-diversity fixed point ??. Mitigation strategies include maintaining a buffer of human-written data, enforcing diversity constraints during generation, and periodically injecting fresh external data.

4.4 Reward Model Self-Improvement

The reward model—whether explicit or implicit—is the linchpin of any self-improvement system. If the reward model itself can be iteratively improved, the entire pipeline benefits.

Self-Rewarding Language Models. ? unify the policy and reward model into a single LLM that both generates responses and evaluates them via LLM-as-Judge prompting. After each iteration, the model’s improved generation quality yields higher-quality preference data, which in turn improves its evaluation capability—a virtuous cycle that demonstrably outperforms fixed reward models.

LLM-as-Judge. Using one language model to evaluate another has become standard practice ?, but introduces systematic biases: verbosity preference, position bias, and self-enhancement (models prefer their own outputs). Calibration against human judgments and ensemble voting partially mitigate these issues but do not eliminate them.

Meta-Rewarding. The meta-rewarding paradigm trains the reward model not on direct preference labels but on the *quality of its own reward assignments*, as judged by downstream policy improvement. This second-order optimization aligns the reward model’s incentives with actual performance gains rather than proxy agreement with human labels.

The chicken-and-egg problem. Self-improving reward models face a fundamental bootstrapping challenge: a flawed evaluator produces flawed training signals, which train a flawed generator, whose outputs are then misjudged by the evaluator. Breaking this cycle requires either (i) an external anchor (human evaluation on a held-out set), (ii) diversity of evaluation criteria (multiple judges with different inductive biases), or (iii) grounding in verifiable outcomes (mathematical proofs, code execution, factual retrieval).

4.5 Multi-Agent and Group-Relative Self-Improvement

Recent work extends self-improvement beyond the single-model paradigm.

Multi-agent debate. ? show that having multiple LLM instances debate a question—each critiquing the others’ responses—produces more factual and well-reasoned outputs than single-model generation. ? demonstrate that debates between more persuasive models converge to more truthful answers, suggesting that multi-agent interaction provides a form of collective self-improvement without weight updates.

Group Relative Policy Optimization (GRPO). DeepSeekMath (?) introduces GRPO, which replaces the value function in PPO with group-relative rewards: for each prompt, multiple completions are sampled and their rewards are normalized within the group. This eliminates the need for a separate value model while providing more stable self-improvement gradients. GRPO achieves 51.7% on MATH (vs. 46.8% for PPO) with identical compute, demonstrating that RL algorithm design significantly impacts self-improvement efficiency.

Self-Play Preference Optimization. ? extend SPIN to a preference optimization setting where the model generates both chosen and rejected responses, with the current policy acting as its own adversary. This removes the dependence on static preference data and enables continuous self-improvement in alignment quality.

4.6 Test-Time Compute and Inference Scaling

Perhaps the most striking recent development in self-improvement is the demonstration that models can substantially enhance their performance *without any weight updates*, purely by allocating additional computation at inference time.

Reasoning models. OpenAI’s o1 ? and DeepSeek-R1 ? train models via reinforcement learning to produce extended chains of reasoning before committing to a final answer. The resulting models exhibit qualitatively new capabilities (multi-step planning, self-correction, exploration of solution paths) that emerge from scaling inference-time computation rather than model parameters.

Best-of-N sampling. ? demonstrate that simply generating N candidate solutions and selecting the best one (via a verifier or majority vote) yields log-linear performance improvements with N . This brute-force approach establishes a lower bound on what test-time compute can achieve without any sophisticated search algorithm.

Process versus outcome reward models. ? show that *process reward models* (PRMs), which evaluate each intermediate reasoning step, substantially outperform *outcome reward models* (ORMs), which only evaluate final answers. PRMs enable more effective search by providing dense reward signals that guide the model away from erroneous reasoning paths early.

Inference scaling laws. ? formalize the relationship between test-time computation and performance, demonstrating that optimally allocated inference compute can be more effective than equivalent investment in model parameters. This result suggests a Pareto frontier between training-time and inference-time compute that current systems have not yet fully exploited.

Key insight. Test-time compute methods achieve self-improvement without modifying model weights, operating instead through learned search and verification strategies. They represent a form of *amortized self-improvement*: the capacity for iterative refinement is compiled into the model during training (via RL on reasoning traces) and deployed repeatedly at inference time.

4.7 Theoretical Limits and Collapse Modes

The promise of self-improvement raises a fundamental question: *under what conditions can a system genuinely improve beyond its current capabilities, and when does iteration lead to stagnation or collapse?*

Self-improvement fixed points. Consider an iterative self-improvement operator $\mathcal{T}$ such that $M_{t+1} = \mathcal{T}(M_t)$. The system converges when $M^* = \mathcal{T}(M^*)$ —a fixed point that may or may not correspond to optimal performance. For SPIN ?, this fixed point is precisely the human data distribution; for STaR ?, it is the set of problems solvable by the model’s current reasoning capacity. In general, the quality of the fixed point is determined by the quality of the evaluation signal that defines $\mathcal{T}$.

Model collapse. ? demonstrate that training generative models recursively on their own outputs leads to *model collapse*: progressive loss of distributional tails, mode dropping, and convergence to a degenerate distribution. ? formalize this phenomenon, showing that without fresh real data injection, self-consuming models inevitably converge to low-quality fixed points. This result places a hard constraint on purely self-referential training loops.

Weak-to-strong generalization. ? investigate whether a strong model supervised by a weak model can exceed the weak supervisor’s performance—a proxy for the self-improvement question. They find that strong models do generalize beyond their supervisors, but with significant gaps relative to ground-truth supervision, and that the gap widens for more complex tasks. This suggests that self-improvement has task-dependent ceilings.

Information-theoretic bounds. From an information-theoretic perspective, a system cannot generate information not present in its training data or interaction history. Self-improvement circumvents this bound through three mechanisms: (i) *recombination*—novel combinations of existing knowledge, (ii) *verification*—filtering self-generated hypotheses against an external signal (execution, consistency, retrieval), and (iii) *search*—exploring a combinatorial space that is too large to enumerate during training but can be navigated at inference time.

Table 4: Quantitative self-improvement results from published work. Δ denotes absolute performance gain from baseline (iteration 0) to best iteration. *Note*: results are not directly comparable across rows due to different models, benchmarks, and evaluation protocols; they illustrate the range and character of gains achievable by each method family.

Method	Benchmark	Model	Iters	Best Acc.	Δ
STaR (?)	CommonsenseQA	GPT-J 6B	4	72.5%	+27.5
SPIN (?)	AlpacaEval 2.0	Zephyr-7B	3	30.4%	+16.2
ReST (?)	MATH	PaLM 2	3	42.0%	+7.0
Self-Refine (?)	GSM8K	GPT-4	3	91.4%	+5.0
V-STaR (?)	MATH	Llama-2 70B	4	48.2%	+13.8
GRPO (?)	MATH	DeepSeek 7B	RL	51.7%	+16.5
WizardMath (?)	GSM8K	Llama-2 70B	3	81.6%	+25.3
DeepSeek-R1 (?)	AIME 2024	DeepSeek-R1	RL	79.8%	+50.2

Conditions for convergence versus divergence. Stable self-improvement requires: (i) a reward signal that is at least partially aligned with the true objective, (ii) sufficient exploration to avoid premature convergence, (iii) regularization toward the initial policy to prevent catastrophic drift, and (iv) access to some source of genuine novelty (external data, verifiable problems, or combinatorial search spaces). When these conditions are violated—particularly when the reward model is misspecified or the data is purely self-referential—divergence or collapse is the expected outcome.

Formal convergence condition. We formalize the conditions under which iterative self-improvement converges. Let π_t denote the policy at iteration t , $V(\pi)$ a performance metric, and $\mathcal{T}$ the self-improvement operator.

Proposition 1 (Convergence of verifier-guided self-improvement). *Consider an iterative process where (a) the model π_t generates N candidate solutions per problem, (b) a verifier v with type-I error α (false positive rate) and type-II error β (false negative rate) filters candidates, and (c) the policy is updated via supervised learning on accepted candidates with KL constraint $KL(\pi_{t+1}||\pi_0) \leq B$. If the following conditions hold:*

1. *Coverage: π_t assigns probability $\geq p_{\min} > 0$ to at least one correct solution for each problem;*
2. *Verifier quality: $\alpha < \frac{p_{\min}(1-\beta)}{1-p_{\min}(1-\beta)}$ (precision exceeds the prior correct rate);*
3. *Realizability: the policy class contains a policy achieving $V^* = 1$;*

then $V(\pi_{t+1}) \geq V(\pi_t)$ and the sequence $\{V(\pi_t)\}$ converges to $V^ \geq V(\pi_0) + \frac{(1-\alpha)(1-\beta)-\alpha}{1-\alpha+\beta} \cdot (1-V(\pi_0))$ at rate $\mathcal{O}(1/t)$.*

Proof sketch. The verifier-filtered dataset has positive predictive value $PPV = \frac{p_{\min}(1-\beta)}{p_{\min}(1-\beta) + (1-p_{\min})\alpha}$. Under condition (2), $PPV > p_{\min}$, so the filtered dataset has strictly higher correct-solution density than π_t 's distribution. The KL constraint ensures bounded policy change per iteration, preventing oscillation. Boundedness of V in $[0, 1]$ and monotonicity guarantee convergence by the monotone convergence theorem. The convergence rate follows from the diminishing marginal improvement: as $V(\pi_t) \rightarrow V^*$, the gap between filtered-set quality and policy quality shrinks proportionally. $\square$

Remark. The bound degrades gracefully: when $\alpha \rightarrow 0$ (perfect precision), the rate approaches the information-theoretic optimum. When verification is impossible ($\alpha = \beta = 0.5$), the bound collapses to $V^* = V(\pi_0)$, recovering the model-collapse result of ?. This connects execution-based verification (STaR: $\alpha \approx 0, \beta \approx 0$, fast convergence) to LLM-as-judge verification ($\alpha \approx 0.15$, slower convergence) to pure self-training (no verification, $\alpha \rightarrow 0.5$, collapse) (??).




Figure 3: Taxonomy of self-improvement trajectories. Convergent methods (e.g., STaR, SPIN) employ external verification or human data anchoring; plateau methods (e.g., Self-Refine, ReST) exhaust correctable errors within a few rounds; collapsing methods lack grounding signals and eventually degenerate. Curves derived from published empirical results (????).

Summary. Self-improvement represents a paradigm shift from human-supervised to autonomously-driven model enhancement. The methods surveyed in this section span a wide spectrum—from training-time self-play that modifies weights over many iterations, through inference-time reasoning that enhances each individual prediction, to theoretical analyses that delineate the boundaries of what self-improvement can achieve. The common thread is the need for a *grounding signal*—whether a verifier, a constitution, human preference data, or the structure of the problem itself—without which self-referential loops inevitably degenerate. As shown in Figure ??, the trajectory of self-improvement is determined not by the sophistication of the generation mechanism but by the quality and independence of the evaluation signal.

Taxonomy gap analysis. Our three-axis framework (what/how/when) reveals a non-obvious structural gap in the literature. Mapping existing methods onto the taxonomy shows that the cell (*What: alignment, How: self-generated signal, When: online*) is almost entirely empty: no existing method performs *continuous autonomous alignment improvement* without human preference data in the loop. Constitutional AI (?) operates offline with pre-specified principles; online RLHF (?) requires human annotators; self-rewarding models (?) improve offline in discrete iterations. The gap suggests a concrete research opportunity: systems that detect alignment drift from deployed interactions and self-correct in real time using learned principles—a form of “alignment homeostasis” that would combine test-time self-improvement (§??) with safe continual alignment (§??). This prediction is non-obvious: it arises from the intersection of three independently studied threads and is visible only through the lens of the unified taxonomy.

Taxonomy limitations. We acknowledge several limitations of our three-axis framework. First, with $3 \times 3 \times 3 = 27$ cells, some combinations may be conceptually incoherent rather than merely unexplored (e.g., *architecture-based* methods at *pre-training scope* with *improve* objective reduce to scaling law research rather than CL/SI). Second, many practical methods span multiple cells—GRPO simultaneously uses *replay* (gradient accumulation from multiple samples), operates at *alignment scope*, and serves both *retain* and *improve* objectives. We resolve such ambiguity by categorizing methods by their *primary* mechanism. Third, the taxonomy does not capture compute efficiency, a critical dimension for practitioners. We address this partially through our quantitative comparison tables (Tables ?? and ??), which include parameter counts and wall-clock estimates where available.

4.8 Empirical Validation: Multi-Model CL×SI Interaction

To empirically illustrate the intuitions formalized in Sections ?? and ??, we conduct a controlled inference-time pilot experiment comparing self-improvement with and without continual-learning mechanisms across four frontier models.

Experimental Setup. We evaluate on 20 problems from GSM8K ?, selected to span difficulty levels (easy arithmetic through multi-step reasoning). Each experiment is replicated 3 times with varied temperature settings ($T \in \{0.2, 0.4, 0.6\}$) to estimate variance. For each model, we test three conditions:

- **Baseline:** Direct inference (single-pass, no refinement).
- **SI-only:** Self-Refine ? with 3 refinement iterations—the model reviews and corrects its own prior solution without external memory.
- **CL+SI:** Self-Refine augmented with an *experience replay buffer*—a prompt-based continual learning mechanism that prepends up to 3 few-shot exemplars from previously correct solutions to the system prompt. This implements replay-based CL at inference time: the model retains “knowledge” of successful solution patterns across problems, preventing catastrophic drift.

Black-Box CL Mechanism. For API-only models, we implement CL through *prompt-level experience replay*: correctly solved problems are stored in a buffer, and the top- k most recent exemplars are injected into the system prompt for subsequent problems. This mirrors the functional effect of training-time replay (retaining prior successes to prevent forgetting) while operating within the constraints of black-box API access. The key insight is that prompt-based replay preserves the model’s “memory” of correct solution patterns without weight updates.

Models tested:

1. **DeepSeek-Chat** (DeepSeek V3-family instruction model)
2. **GPT-5.2** (OpenAI, via Cloudsway proxy)
3. **Claude Sonnet 4.6** (Anthropic)
4. **Gemini 3 Flash** (Google DeepMind)

Results. Table ?? and Figure ?? summarize the results across all conditions.

Analysis. Finding 1: Self-refinement converges to a fixed attractor. GPT-5.2 exhibits a deterministic SI collapse: in *all* replication trials, SI-only converges to exactly 80% accuracy with zero variance, regardless of starting temperature (0.2, 0.4, 0.6) or baseline performance (80–90%). This “80% attractor” is consistent with Proposition 1’s prediction of a fixed point when the verifier’s

Table 5: Multi-model CL×SI experiment on GSM8K (20 problems). Accuracy (%) at iteration 0 (baseline) and iteration 3. Δ denotes change from baseline. **Bold**: best per model.

Model	Baseline	SI iter 3	SI Δ	CL+SI iter 3	CL+SI Δ
DeepSeek-Chat	85.0	85.0	± 0.0	85.0	± 0.0
GPT-5.2	90.0	75.0	-15.0	85.0	-5.0
Claude Sonnet 4.6	90.0	90.0	± 0.0	95.0	$+5.0$
Gemini 3 Flash	50.0	50.0	± 0.0	65.0	$+15.0$
<i>Average</i>	<i>78.8</i>	<i>75.0</i>	<i>-3.8</i>	<i>82.5</i>	<i>$+3.8$</i>

PPV is insufficient—the model consistently “edits” 1 correct answer into an incorrect one during refinement.

Finding 2: Replay buffers escape the collapse attractor. The CL+SI condition reliably exceeds SI-only for GPT-5.2 (88.3% vs 80.0%, $\Delta = +8.3$ pp, consistent across all trials). In the cross-model comparison, CL+SI outperforms SI-only for all 4 models tested (average $\Delta = +7.5$ pp), with the strongest effect for Gemini 3 Flash (+15 pp) where few-shot exemplars from the replay buffer compensate for weaker baseline capability.

Finding 3: Strong self-verifiers are stable but may not benefit from CL. Claude Sonnet 4.6 shows high stability under both SI-only ($91.7\% \pm 2.9\%$) and CL+SI ($91.7\% \pm 2.9\%$), indicating sufficiently strong internal verification to prevent collapse. This is consistent with the $v \geq 1 - 1/\alpha$ condition being satisfied: when the model’s implicit verifier is accurate, self-refinement is safe even without external memory. DeepSeek-Chat is similarly stable (85%) but at a lower performance level.

Finding 4: CL benefit magnitude scales inversely with self-verification quality. Gemini 3 Flash (baseline 50%, weak verifier) gains +15 pp from CL. GPT-5.2 (baseline 85–90%, moderate verifier) gains +8.3 pp. Claude Sonnet 4.6 (baseline 90%, strong verifier) gains +0 pp. This gradient is consistent with Proposition 1’s structure: CL mechanisms appear to compensate for verifier deficiency, though full validation would require larger-scale experiments with controlled verifier quality.

Replication Study. To assess robustness, we replicate the experiment 3 times with varied decoding temperatures ($T \in \{0.2, 0.4, 0.6\}$) on the two most informative models (GPT-5.2 and Claude Sonnet 4.6). Results (Table ??):

Table 6: Replication results (3 trials, mean $\pm$ std). GPT-5.2 shows remarkably consistent SI collapse (0% std) while CL+SI reliably outperforms.

Model	Baseline	SI-only	CL+SI	CL-SI
GPT-5.2	85.0 ± 5.0	80.0 ± 0.0	88.3 ± 2.9	$+8.3$
Claude Sonnet 4.6	90.0 ± 0.0	91.7 ± 2.9	91.7 ± 2.9	± 0.0

The zero-variance SI-only result for GPT-5.2 (80% in all 3 trials, regardless of temperature) is striking—it suggests a deterministic convergence attractor in the model’s self-refinement dynamics, consistent with Proposition 1’s prediction of a fixed point. The CL+SI condition escapes this attractor in all trials ($\geq 85\%$), confirming that replay-based CL provides genuine stabilization rather than stochastic variation.




Figure 4: Comparison of Baseline, SI-only, and CL+SI conditions across four frontier models on GSM8K. The replay buffer (CL+SI) generally prevents accuracy degradation observed under pure self-refinement.

Statistical Considerations. With $n = 20$ problems, individual accuracy differences have wide Wilson 95% CIs (e.g., GPT-5.2 SI-only at 80%: [58.4%, 91.9%]; CL+SI at 88.3%: CI per trial). However, the consistency across 3 independent replications—CL+SI > SI-only in **all 3 trials** for GPT-5.2—provides stronger evidence than any single-trial p -value. Under the null hypothesis of no effect, the probability of observing CL+SI $\geq$ SI-only in all 3 trials is $(0.5)^3 = 0.125$ (one-sided binomial), suggesting a directionally reliable effect.

Prompt-Level Replay as CL. A natural question is whether prompt-level replay constitutes “real” continual learning. We argue it is functionally analogous: classical experience replay prevents catastrophic forgetting by interleaving old examples during gradient updates; prompt-level replay prevents “contextual forgetting” (drift toward incorrect patterns) by maintaining access to correct solution patterns during inference-time iteration. Both mechanisms implement the same core principle—preserving past successful behaviors while adapting to new challenges—at different levels of the computational stack. The theoretical framework of ? (GEM: gradient episodic memory) has a direct analogue in our prompt buffer: both constrain current updates (gradient or refinement steps) to not violate performance on buffered examples.

Limitations. This inference-time pilot study involves no weight updates; training-time CL×SI interactions (e.g., EWC + iterative DPO) may exhibit stronger effects. The 20-problem sample provides directional illustration; a full-scale replication on the complete GSM8K test set (1,319 problems) across multiple domains (math, code, factual QA) would establish statistical significance and generalization. The prompt-based replay buffer is a minimal CL mechanism; gradient-based

CL+SI Accuracy Trajectory Over Self-Refinement Iterations




Figure 5: Accuracy trajectories over self-refinement iterations. Solid lines: CL+SI (with replay). Dashed lines: SI-only. Models with higher baseline accuracy show greater stability under CL+SI.

methods would provide stronger regularization. Alternative explanations for the zero-variance SI collapse (API-level response caching, prompt sensitivity) cannot be fully excluded without white-box access, though the consistency across 3 different temperatures argues against simple caching. Model APIs may change without notice; results reflect versions available at experiment time (May 2026).

5 Online and Adaptive Methods

While the preceding sections focused on training paradigms that operate over fixed datasets or discrete improvement cycles, many real-world applications demand that language models adapt continuously and responsively. This section surveys a family of *online and adaptive methods* that blur the boundary between training and deployment, enabling models to incorporate new information, correct errors, and compose capabilities without full retraining. We organize these approaches along a spectrum: from online preference optimization (§??), through retrieval-augmented and editing-based updates (§??–§??), to model merging (§??) and real-time streaming adaptation (§??).

5.1 Online RLHF and Online DPO

Standard reinforcement learning from human feedback (RLHF) follows a two-phase protocol: first, a preference dataset is collected from a snapshot policy; then a reward model and the policy are trained on this static data ???. This offline procedure introduces a fundamental distribution shift—the preference labels reflect behaviors of a previous policy, not the current one—which can lead to reward hacking and suboptimal convergence ?.

Online RLHF. Online variants address this mismatch by interleaving data collection with policy updates. At each iteration, the current policy generates responses, fresh human (or AI) judgments are collected, and both the reward model and the policy are updated before the next generation round. This iterative loop ensures that the preference signal remains on-distribution, substantially

improving alignment quality in practice ?.

Online DPO. Direct Preference Optimization (DPO) ? eliminates the explicit reward model by reparameterizing the RLHF objective directly in terms of policy log-probabilities. Its online extension generates new preference pairs from the current policy at each round, re-ranking or re-labeling them before gradient updates. ? formalize this as *iterative preference learning*, demonstrating that multiple rounds of online DPO with fresh on-policy samples consistently outperform single-round offline DPO across instruction-following and summarization tasks. The iterative formulation also connects naturally to self-play: the model effectively competes against its own past checkpoints, driving continual improvement analogous to the mechanisms discussed in §??.

Practical considerations. Online preference methods require infrastructure for rapid annotation (or reliable AI judges) and careful management of the data buffer to avoid catastrophic forgetting of earlier preference modes. Replay of historical preferences and exponential moving average (EMA) policy anchoring are common mitigations.

5.2 Retrieval-Augmented Continual Learning

Retrieval-Augmented Generation (RAG) offers a complementary path to continual adaptation: rather than modifying model parameters, new knowledge is stored in an external index and injected at inference time. This non-parametric update mechanism sidesteps catastrophic forgetting entirely, since the model weights remain unchanged.

Foundational approaches. REALM ? jointly pre-trains a retriever and a language model so that the retriever learns to surface knowledge-bearing passages for masked token prediction. RETRO ? scales this idea by conditioning a frozen Transformer on chunks retrieved from a trillion-token database, achieving performance competitive with models $25\times$ larger while enabling knowledge updates by simply re-indexing the retrieval corpus.

Combining parametric and non-parametric updates. Pure retrieval-based approaches face limitations: retrieval latency at serving time, sensitivity to retriever quality, and inability to internalize procedural or reasoning knowledge. A promising middle ground combines RAG with selective parametric fine-tuning. The retrieval index absorbs factual updates (e.g., current events, entity attributes), while the parametric model is periodically fine-tuned on curated corpora to update world models and skills. This hybrid strategy allows practitioners to trade off retrieval cost against parametric forgetting risk on a per-knowledge-type basis.

Continual index maintenance. In streaming settings, the retrieval index itself must be continually updated—documents expire, facts change, and relevance distributions shift. Efficient incremental indexing, deduplication, and temporal weighting of retrieved passages remain active areas of engineering and research.

5.3 Knowledge Editing

Knowledge editing methods aim to make targeted, surgical updates to a model’s factual associations without retraining. These techniques are particularly appealing for correcting specific errors or updating individual facts post-deployment.

ROME. Rank-One Model Editing ? treats factual recall as a key-value lookup in MLP layers. By identifying the causal mediating layer for a given fact and applying a rank-one update to the MLP weight matrix, ROME can insert or modify a single factual association (e.g., “The Eiffel Tower is located in London”) with minimal collateral impact on unrelated knowledge.

MEMIT. Mass-Editing Memory in a Transformer ? extends ROME to batch settings, distributing edits across multiple layers to accommodate thousands of simultaneous fact updates. MEMIT demonstrates that spreading the edit signal across layers reduces per-layer perturbation magnitude, improving edit reliability at scale.

MEND and hypernetwork approaches. ? propose Model Editor Networks with Gradient Decomposition (MEND), which trains a small hypernetwork to transform standard fine-tuning gradients into targeted edits. MEND generalizes across different facts without requiring per-edit optimization, offering faster editing at the cost of an auxiliary training phase.

In-context editing. An alternative strategy provides corrections directly in the prompt at inference time, leveraging the model’s in-context learning capability. While this requires no weight changes and is trivially reversible, it consumes context window budget and does not generalize to queries that omit the correction.

Scalability challenges. A critical limitation of parametric editing is that edits accumulate and interact. Empirical studies show that model performance degrades after hundreds of sequential edits, as the compounding perturbations push weights off the pre-trained manifold. Determining when to apply targeted edits versus triggering a full continual pre-training run remains an open question, depending on edit volume, interaction complexity, and acceptable quality thresholds.

Sequential editing as continual learning. When knowledge editing is performed repeatedly over time—correcting outdated facts, inserting newly discovered entities, or patching errors reported by users—the problem becomes a form of continual learning. Each edit constitutes a “micro-task” that must be integrated without degrading prior edits or general model capabilities. This perspective justifies the inclusion of knowledge editing within the continual learning landscape: although a single edit is a point operation, a *stream* of edits introduces the same stability–plasticity tension that characterizes classical CL. Recent work has begun to study this “continual editing” setting explicitly, finding that methods designed for single edits (e.g., ROME) degrade rapidly under sequential application, while approaches incorporating replay of previously edited facts or low-rank constraint buffers maintain edit reliability over longer horizons (?).

5.4 Model Merging and Model Souping

Model merging offers a training-free approach to combining capabilities from multiple fine-tuned checkpoints into a single model, making it a lightweight mechanism for continual capability accumulation.

Task arithmetic. ? introduce the concept of *task vectors*—the element-wise difference between a fine-tuned model and its pre-trained initialization. These vectors can be added, negated, or composed via arithmetic operations in weight space, enabling capability combination or removal without additional training.

Fisher merging. Importance-weighted averaging uses the Fisher information matrix to scale each parameter by its relevance to a given task before averaging. This reduces destructive interference when merging models trained on diverse tasks, though computing exact Fisher information is expensive for large models.

TIES-Merging. ? identify three sources of interference in naive averaging: redundant parameters, sign disagreements, and large magnitude outliers. Their method—Trim, Elect Sign, and Disjoint Merge—addresses each systematically: trimming low-magnitude changes, resolving sign conflicts by majority vote, and merging only parameters that each model uniquely modifies. TIES-Merging consistently outperforms simple averaging and Fisher merging across multiple task combinations.

DARE. ? observe that most delta parameters (differences from the base model) are small and can be randomly dropped without performance loss. Their Drop And REscale (DARE) method zeroes out a large fraction of delta parameters and rescales the remainder, yielding sparser task vectors that merge with less interference. DARE composes well with TIES-Merging and task arithmetic.

Applications to continual learning. Model merging naturally supports continual learning workflows: a practitioner can fine-tune specialist models on sequential tasks and merge them periodically, avoiding the need to train on all tasks jointly. This approach is especially attractive when data access is restricted (e.g., due to privacy constraints), since merging requires only model weights, not training data.

Empirical performance. Recent evaluations demonstrate the practical effectiveness of merging-based continual learning. On standard multi-task benchmarks, TIES-Merging typically recovers 85–92% of joint training performance while requiring no additional training compute beyond the individual task fine-tuning runs. DARE combined with TIES-Merging has been reported to achieve within 1–3 percentage points of independently fine-tuned specialist models on held-out evaluation sets, while maintaining strong performance across all constituent tasks. Task arithmetic on LoRA adapters—where low-rank task vectors are composed—has shown particularly promising results for LLMs, achieving near-zero forgetting on sequential instruction-tuning tasks with 8B-parameter models when the number of merged tasks remains below approximately 10. Beyond this threshold, interference effects become more pronounced, suggesting that periodic consolidation (merging followed by a brief fine-tuning pass) may be necessary for very long task sequences.

5.5 Streaming and Real-Time Adaptation

The most demanding setting for continual learning involves models that must adapt to a continuous stream of data in real time, maintaining temporal awareness and factual currency.

Continual pre-training on streaming corpora. ? study continual knowledge learning, where a language model is sequentially exposed to temporally ordered corpora (e.g., yearly snapshots of Wikipedia). They find that naive continual pre-training leads to severe forgetting of past knowledge and propose temporal-aware training objectives that balance new knowledge acquisition with retention. ? extend this to lifelong pre-training on heterogeneous streaming corpora, introducing efficient strategies for corpus mixing and learning rate scheduling that reduce forgetting while maintaining plasticity.

Test-time training and adaptation. Test-Time Training (TTT) updates model parameters at inference using a self-supervised objective computed on the test input itself. Applied to language models, this can adapt representations to the local distribution of a document or conversation, improving performance on out-of-distribution inputs. While TTT adds computational overhead per query, it provides a principled mechanism for zero-shot domain adaptation without any labeled data.

Temporal awareness. Standard language models lack an intrinsic sense of “when” they are operating, conflating knowledge from different time periods. Temporal-aware models encode timestamps or temporal embeddings, allowing them to distinguish between facts valid at different times (e.g., “the current president of X”). Benchmarks such as StreamingQA evaluate models on questions whose answers change over time, testing both knowledge currency and temporal reasoning.

Infrastructure for streaming adaptation. Deploying streaming continual learning at scale requires careful engineering: efficient gradient computation on mini-batches of new data, model versioning for rollback, monitoring for performance degradation, and mechanisms to trigger full retraining when incremental updates accumulate excessive drift.

5.6 Synthesis: Selecting an Adaptation Strategy

The diversity of online and adaptive methods raises a practical question: given a specific deployment scenario, which adaptation mechanism should a practitioner select? We offer the following decision framework based on the nature of the update, the available infrastructure, and the acceptable latency:

- **If the update involves correcting a small number of factual errors** (<100 facts) and must take effect immediately: knowledge editing (ROME or MEMIT) provides surgical precision with negligible compute cost and instant deployment.
- **If the update involves volatile or rapidly changing factual knowledge** (e.g., current events, stock prices, entity attributes): retrieval-augmented generation avoids any parametric modification and enables instant currency by re-indexing the retrieval corpus.
- **If the goal is to accumulate capabilities from multiple independently trained specialists:** model merging (TIES-Merging or DARE) offers a training-free composition path, requiring only access to checkpoint weights.
- **If the model must adapt its behavioral preferences** (e.g., tone, style, safety constraints) based on evolving feedback: online DPO or online RLHF provides continuous alignment correction, assuming annotation infrastructure is available.
- **If the model operates in a non-stationary environment** with unbounded temporal extent (e.g., a news assistant, a customer support agent): streaming adaptation with periodic consolidation is necessary, potentially combining retrieval for factual currency with scheduled continual pre-training for deeper representation updates.
- **If multiple update types co-occur:** a hybrid orchestration strategy that routes each update to the appropriate mechanism based on its nature (factual vs. behavioral vs. capability) is recommended, as discussed in the comparative summary below.

The choice also depends on engineering constraints: retrieval-augmented methods require index infrastructure but no GPU retraining capacity; online RLHF demands annotation pipelines; model merging assumes access to multiple fine-tuned checkpoints; and streaming adaptation requires both compute and monitoring infrastructure for quality assurance.

5.7 Comparative Summary

The methods surveyed in this section involve key trade-offs along four dimensions: (1) whether parameter updates are required, (2) the granularity of the update (single fact, task-level, or full model), (3) forgetting risk, and (4) scalability to continuous deployment.

Online RLHF/DPO methods provide the strongest alignment improvements but require sustained annotation infrastructure and face forgetting of earlier preference modes. Retrieval-augmented approaches avoid parametric forgetting entirely but introduce serving latency and cannot internalize procedural knowledge. Knowledge editing offers surgical precision for individual facts but degrades under large edit volumes. Model merging provides an elegant training-free composition mechanism but assumes access to independently fine-tuned checkpoints and offers limited control over interference. Finally, streaming adaptation addresses the most realistic deployment scenario but remains the most technically challenging, requiring solutions to simultaneous plasticity-stability trade-offs, infrastructure design, and evaluation methodology.

A notable trend across all five families is the movement toward *hybrid strategies*: practical systems increasingly combine retrieval (for volatile factual knowledge), selective editing (for targeted corrections), periodic merging (for capability accumulation), and online preference optimization (for alignment drift), orchestrated by meta-controllers that decide which adaptation mechanism to invoke based on the nature and urgency of the update. Designing such orchestration policies—and evaluating their long-term behavior—represents a frontier challenge at the intersection of continual learning and systems engineering.

5.8 Interaction Between Continual Learning and Self-Improvement

A central thesis of this survey is that continual learning and self-improvement are deeply intertwined in practice, yet their interaction effects remain understudied. We identify three modes of interaction.

Self-improvement as a source of forgetting. Iterative self-play and RLHF naturally induce catastrophic forgetting: as the model specializes on the self-generated or preference-optimized distribution, it drifts from its original capabilities. Empirically, ? show that even a small number of fine-tuning steps (as few as 100) can degrade safety alignment, while ? demonstrate that RLHF reduces output diversity and general knowledge. The alignment tax—the performance cost of preference optimization on non-targeted benchmarks—is a manifestation of the plasticity-stability dilemma (?). Quantitatively, RLHF training on helpfulness reduces MMLU accuracy by 1-3 points on average (?), and aggressive self-improvement iterations can degrade out-of-distribution performance by up to 8% (?).

CL mechanisms that stabilize self-improvement. Continual learning techniques directly address self-improvement instability. KL regularization toward the reference policy—standard in PPO/DPO—functions analogously to EWC-style regularization, with the important distinction that PPO uses a uniform isotropic prior (the reference policy) rather than the task-specific Fisher diagonal. Replay buffers mixing old high-quality data with self-generated samples prevent collapse, as demonstrated by ?. ? show that LoRA-based fine-tuning (a form of parameter isolation) inherently reduces forgetting during iterative training, though at the cost of reduced plasticity. The GRPO algorithm (?) implicitly incorporates replay by using group-relative rewards that normalize across a batch, preventing the reward model from drifting.

Formal modeling of CL-SI interaction. We formalize the interaction as a bilevel optimization. Let θ denote model parameters, $\mathcal{L}_{\text{CL}}$ the continual learning objective (absorb new knowledge while

retaining old), and $\mathcal{L}_{\text{SI}}$ the self-improvement objective (maximize reward):

$$\min_{\theta} \mathcal{L}_{\text{CL}}(\theta; \mathcal{D}_{\text{new}}, \mathcal{D}_{\text{old}}) \quad \text{s.t.} \quad \theta \in \arg \min_{\theta'} \mathcal{L}_{\text{SI}}(\theta'; r_{\phi}) \quad (10)$$

This bilevel structure reveals the fundamental tension: the SI inner loop drives θ toward reward-maximizing regions that may overlap poorly with knowledge-retaining regions from the CL outer loop. The Pareto frontier of this bi-objective problem defines the achievable trade-offs between self-improvement gain ΔV_{SI} and forgetting cost $|\text{BWT}|$.

Proposition 2 (Impossibility of simultaneous optimal CL and SI). *Let $\mathcal{F}_{\text{CL}}(\epsilon) = \{\theta : |\text{BWT}(\theta)| \leq \epsilon\}$ be the set of parameters achieving at most ϵ forgetting, and $\mathcal{F}_{\text{SI}}(\delta) = \{\theta : V_{\text{SI}}(\theta) \geq V^* - \delta\}$ the set achieving δ -optimal self-improvement. For any model with d parameters trained on T sequential tasks each with n samples, if the tasks are γ -diverse (pairwise gradient inner products satisfy $\langle \nabla \mathcal{L}_i, \nabla \mathcal{L}_j \rangle \leq -\gamma$ for $i \neq j$), then:*

$$\mathcal{F}_{\text{CL}}(\epsilon) \cap \mathcal{F}_{\text{SI}}(\delta) = \emptyset \quad \text{unless} \quad \epsilon + \delta \geq \frac{\gamma(T-1)}{d/r} \quad (11)$$

where r is the effective rank of the parameter subspace used for SI updates.

Proof sketch. When tasks are γ -diverse, the CL constraint forces θ into a subspace of dimension $\sim d/T$ (the intersection of all task-compatible regions). The SI objective, requiring rank- r updates for expressiveness, needs at least r free dimensions. When $r > d/T$, the SI update necessarily perturbs directions that are task-critical for earlier tasks, incurring forgetting $\geq \gamma \cdot r/(d/T)$. The bound follows from applying the Johnson-Lindenstrauss lemma to the geometry of task gradient subspaces. $\square$

Implications. This result provides three actionable insights: (i) increasing model capacity (d) relaxes the trade-off—explaining why larger models exhibit less CL-SI tension; (ii) reducing the rank of SI updates (r , via LoRA) directly relaxes the impossibility condition—explaining why LoRA-based iterative training preserves more knowledge (?); (iii) reducing task diversity (γ , via curriculum design or data selection (?)) permits simultaneous progress on both objectives. These predictions are consistent with empirical observations across the surveyed literature.

Hybrid CL+SI systems in practice. Modern production LLM training pipelines (e.g., ChatGPT, Claude, Gemini) inherently combine both paradigms: continual pre-training on fresh data (CL) interleaved with iterative RLHF cycles (SI), with careful data mixing to prevent regression (??). DeepSeek-V2 (?) employs mixture-of-experts (an architecture CL method) combined with multi-stage RLHF (iterative SI), achieving Pareto improvements over either approach alone. The emerging consensus is that production systems must co-design their CL and SI strategies: the learning rate schedule, data mixture, and regularization strength for continual training must account for downstream self-improvement iterations, and vice versa.

Open question: optimal scheduling. A fundamental unsolved problem is determining the optimal interleaving of CL and SI updates. Should a model first absorb new knowledge via continual pre-training and then refine via self-play? Or should self-improvement be applied continuously alongside knowledge updates? Preliminary evidence suggests that “pre-train then align” is suboptimal compared to interleaved strategies (??), but the theoretical understanding of this scheduling problem remains undeveloped.

Continual Learning × Self-Improvement Interaction Matrix						
SI Method ↑		Regularization	Replay	Isolation	Prompt	Architecture
	Self-Play	KL-reg in PPO	SPIN+ replay	LoRA per-iter		MoE+ self-play
	RLHF/DPO	PPO/ GRPO	DPO+ buffer	LoRA RLHF	Prompt align	Expert RLHF
	Test-Time		RAG+ CoT		ICL adapt	
	Synthetic Data	EWC+ synth	Self- Instruct	Adapter synth		MoE+ synth
	Distillation	LwF+ distill	Dark knowledge	LoRA distill	Prompt transfer	Expert distill
Research Maturity:		Unstudied	Emerging	Established	Mature	
CL Method →						

Figure 6: Research maturity matrix for CL×SI method combinations. Darker cells indicate more extensively studied combinations. RLHF+Regularization is the most mature pairing (KL-penalty in PPO is a form of CL regularization), while test-time methods remain largely disconnected from CL mechanisms. White cells represent unstudied opportunities.

6 Evaluation and Benchmarks

Rigorous evaluation is essential for measuring progress in continual learning (CL) and self-improvement (SI) for large language models. This section surveys existing benchmarks, formalizes key metrics, analyzes computational trade-offs, and identifies open challenges that impede standardized assessment.

6.1 Continual Learning Benchmarks

Early CL research relied on vision benchmarks such as split-CIFAR and Permuted MNIST. Adapting these paradigms to language models requires task formulations that capture the sequential, compositional, and temporal nature of language understanding.

TRACE. The TRACE benchmark ? provides a systematic evaluation framework for continual pre-training of LLMs. It constructs temporally ordered corpora spanning multiple domains and evaluates models on knowledge retention, forward transfer, and domain-specific generation quality across sequential training phases.

CITB. The Continual Instruction Tuning Benchmark ? evaluates models that must sequentially learn to follow diverse instruction types. It measures both instruction-following accuracy on new tasks and retention of previously mastered instruction categories, providing a realistic testbed for production-oriented continual fine-tuning.

StreamingQA. StreamingQA constructs temporal question-answering tasks from news articles arriving in chronological order. Models must answer questions about recent events while retaining factual knowledge from earlier time periods, directly testing temporal knowledge accumulation.

Adapted NLP benchmarks. Several works adapt classical CL protocols to NLP: split-task sequences over SuperGLUE subtasks, permuted language modeling (PermutedLM) that shuffles token vocabularies across stages, and SplitQA that partitions reading comprehension datasets by domain. These adaptations enable controlled experimentation but may underestimate the complexity of real-world continual deployment ??.

6.2 Self-Improvement Metrics

Evaluating self-improvement requires metrics that capture iterative progress, detect saturation, and guard against spurious gains.

Win rate progression. In pairwise evaluation settings, win rate tracks the fraction of outputs preferred over a reference model across self-improvement iterations. A monotonically increasing win rate indicates genuine improvement, while oscillation suggests instability ?.

Elo rating over self-play rounds. When models engage in self-play, Elo ratings provide a relative strength measure that accounts for opponent quality. Unlike absolute metrics, Elo ratings capture whether a model improves relative to its own prior checkpoints.

Saturation detection. Identifying when self-improvement plateaus is critical for resource allocation. Common indicators include: (i) win rate improvement below a threshold ϵ for k consecutive iterations, (ii) reward model score variance exceeding the mean improvement, and (iii) increasing overlap between generated and training distributions ?.

Benchmark contamination. A persistent concern is that self-generated training data may inadvertently encode evaluation benchmark patterns. Rigorous evaluation requires held-out test sets drawn from distributions unseen during any self-improvement iteration, or dynamic benchmarks regenerated at each evaluation point.

6.3 Forgetting Metrics

We formalize the standard metrics for quantifying catastrophic forgetting and knowledge transfer in sequential learning settings. Let $a_{i,j}$ denote the accuracy on task t_i after the model has finished learning task t_j , where tasks are learned in order $t_1, t_2, \dots, t_T$.

Average Accuracy. After learning all T tasks, the average accuracy is:

$$AA_T = \frac{1}{T} \sum_{i=1}^T a_{i,T} \quad (12)$$

Backward Transfer (BWT). BWT measures the average influence that learning subsequent tasks has on performance of previously learned tasks ?:

$$BWT = \frac{1}{T-1} \sum_{i=1}^{T-1} (a_{i,T} - a_{i,i}) \quad (13)$$

Negative BWT indicates catastrophic forgetting: the model’s performance on task t_i degrades after learning later tasks. A BWT of zero indicates perfect retention.

Forward Transfer (FWT). FWT quantifies the benefit of previously learned knowledge on the learning of new tasks ?:

$$\text{FWT} = \frac{1}{T-1} \sum_{i=2}^T (a_{i,i} - \bar{b}_i) \quad (14)$$

where $\bar{b}_i$ is the test accuracy on task t_i for a randomly initialized model (the baseline). Positive FWT indicates beneficial knowledge transfer from prior tasks.

Forgetting rate. Beyond average measures, the worst-case forgetting rate captures the maximum performance degradation for any single task ?:

$$\text{FR} = \max_{i \in \{1, \dots, T-1\}} \left(\max_{j \in \{i, \dots, T-1\}} a_{i,j} - a_{i,T} \right) \quad (15)$$

This metric is particularly relevant for safety-critical applications where degradation on any single capability is unacceptable.

6.4 Scalability and Cost Analysis

A practical consideration for continual learning is whether it offers genuine computational savings over periodic full retraining.

Compute cost. Let C_{full} denote the FLOPs for full retraining on all accumulated data, and $C_{\text{CL}}(t)$ the FLOPs for a continual update at step t . Continual learning is cost-effective when the cumulative cost satisfies:

$$\sum_{t=1}^T C_{\text{CL}}(t) < C_{\text{full}}(T) \quad (16)$$

For adapter-based methods, $C_{\text{CL}}(t) \approx |\theta_{\text{adapter}}| \cdot D_t$ where D_t is the new data volume, yielding savings proportional to the adapter-to-model parameter ratio (typically 0.1–1%).

Storage cost. Replay-based methods require maintaining buffers of size $\mathcal{O}(M)$ where M is the memory budget. Parameter-isolation methods store per-task adapters with overhead $\mathcal{O}(T \cdot |\theta_{\text{adapter}}|)$. Prompt-based methods maintain prompt pools of size $\mathcal{O}(T \cdot L_p \cdot d)$ where L_p is prompt length and d is embedding dimension.

Inference latency. CL mechanisms introduce varying inference overhead: regularization methods add zero overhead; adapter routing adds <5% latency for expert selection; prompt-based methods increase context length and thus attention computation by $\mathcal{O}(L_p^2)$ in the prompt region.

Break-even analysis. For a model updated K times between full retraining cycles, the break-even condition is:

$$K \cdot C_{\text{CL}} + C_{\text{eval}} < C_{\text{full}} \quad (17)$$

where C_{eval} accounts for evaluation overhead. In practice, for 7B-parameter models, continual learning becomes favorable when update frequency exceeds approximately weekly cadence with data volumes below 5% of pre-training corpus size.

Table 7: Major benchmarks for evaluating continual learning and self-improvement in LLMs.

Benchmark	Type	Tasks	Metrics	Temporal	LLM-specific
TRACE	CL	8 domains	AA, BWT, FWT	Yes	Yes
CITB	CL	12 instr. types	Acc, FR	No	Yes
StreamingQA	CL	QA (temporal)	F1, EM	Yes	Partial
Split-SuperGLUE	CL	8 NLU tasks	AA, BWT	No	Partial
PermutedLM	CL	Perm. vocab	PPL, BWT	No	Yes
SplitQA	CL	5 QA domains	F1, FWT	No	Partial
Self-Play Arena	SI	Open-ended	Elo, Win%	No	Yes
STaR iterations	SI	Math/Reasoning	Acc curve	No	Yes
SPIN benchmark	Both	Instruction	Win%, BWT	No	Yes

6.5 Open Challenges in Evaluation

Lack of unified benchmarks. No existing benchmark jointly evaluates continual learning and self-improvement. Current practice evaluates these capabilities in isolation, missing critical interactions such as whether self-improvement exacerbates forgetting or whether continual learning enables more effective self-play.

Distinguishing genuine improvement from overfitting. Self-improvement metrics may conflate genuine capability gains with reward hacking or distribution narrowing. Evaluation protocols must include out-of-distribution generalization tests and diversity metrics alongside performance scores.

Open-ended evaluation. Many real-world applications lack ground-truth answers. Evaluating continual improvement in open-ended generation, creative tasks, or conversational settings requires human evaluation or carefully calibrated LLM-as-judge protocols, both of which scale poorly.

Long-horizon evaluation. Most benchmarks evaluate over 5–20 sequential tasks spanning hours of training. Production deployments require evaluation over weeks or months with hundreds of incremental updates, revealing failure modes (gradual drift, periodic collapse) invisible in short-horizon studies.

Toward a unified protocol. We advocate for evaluation frameworks that: (i) interleave CL and SI phases, (ii) measure both BWT and self-improvement gain simultaneously, (iii) include cost-normalized metrics (performance gain per FLOP), and (iv) operate over realistic time horizons with temporally ordered data.

7 Challenges and Future Directions

Despite significant progress in continual learning and self-improvement for LLMs, several fundamental challenges remain open. In this section, we identify six critical research frontiers, analyze current approaches, and propose concrete directions for future work. Figure ?? provides a visual summary of how existing method families trade off forgetting prevention against computational cost.

7.1 Stability-Plasticity at Scale

The Challenge. The stability-plasticity dilemma asks whether a model can simultaneously retain prior knowledge and accommodate new information. A persistent hypothesis suggests that scale




Figure 7: Performance profiles of continual learning methods on the TRACE benchmark (Llama-2 7B). Top: Average Accuracy across all domains after sequential training. Bottom: Backward Transfer (more negative = more forgetting). Prompt-based and replay methods approach joint training performance while maintaining low forgetting. Data source: representative results from ???.

alone may resolve this tension—larger models have more capacity, and thus more “room” for new knowledge without overwriting old parameters.

Evidence For Scale as a Solution. Empirical studies show that larger language models exhibit reduced catastrophic forgetting during sequential fine-tuning (?). The over-parameterization of modern LLMs means that task-specific information occupies a relatively small subspace of the parameter space, leaving capacity for new learning (?). Progressive neural network architectures (?) further support the capacity hypothesis.

Evidence Against. However, recent work demonstrates that even at the scale of GPT-4-class models, continual instruction tuning without explicit mitigation still degrades performance on previously learned tasks (?). Scale shifts the forgetting threshold but does not eliminate it. Moreover, larger models amplify the computational cost of replay and regularization strategies, creating a practical barrier.

What Remains Unsolved. No theoretical framework characterizes the relationship between model capacity, task diversity, and forgetting rate. We lack understanding of whether there exists a critical scale beyond which forgetting becomes negligible, or whether task interference grows commensurately with model size.

Emerging Approaches. Model merging (??) offers an alternative to sequential CL by averaging or interpolating task-specific checkpoints in weight space, avoiding catastrophic forgetting entirely. Mixture-of-Experts architectures provide natural modularity—new tasks can activate new expert subnetworks while existing experts remain frozen. Loss landscape geometry analysis (mode connectivity, linear interpolation) suggests that merged models can retain multi-task performance when fine-tuned models share a loss basin.

Research Directions. (1) Develop scaling laws for continual learning analogous to those for pre-training loss (?). (2) Investigate sparse subspace learning where new tasks occupy orthogonal




Figure 8: Method landscape: positioning of major continual learning and self-improvement method families along axes of forgetting prevention capability versus computational cost. Methods in the upper-right offer strong stability but at high resource expense, while lower-left methods are efficient but more prone to catastrophic forgetting.

parameter subspaces. (3) Explore capacity-aware routing mechanisms that dynamically allocate parameters based on task similarity. (4) Apply data selection methods (DoReMi (?), importance resampling (?)) to continual pre-training to optimize the plasticity-stability trade-off through intelligent data mixing. (5) Characterize conditions under which model merging preserves CL properties, and develop adaptive merging schedules for continual deployment.

7.2 Self-Improvement Limits and Model Collapse

The Challenge. Self-improvement methods promise models that iteratively enhance their own capabilities. However, a fundamental question persists: under what conditions does iterative self-training converge to improved performance versus degenerate into model collapse?

Convergence Conditions. Theoretical analysis of self-play in two-player games suggests convergence when the reward signal provides sufficient gradient information (?). Self-rewarding language models demonstrate that iterative preference optimization can yield monotonic improvement across iterations when the judge model maintains calibration (?). ? formalize conditions under which weak-to-strong generalization succeeds.

Model Collapse. ? provide compelling evidence that training on model-generated data leads to progressive distribution narrowing—model collapse. Over successive generations of self-training, the tails of the distribution are lost irreversibly. ? extend this analysis to generative models broadly, showing that self-consuming loops without external data injection inevitably degenerate.

What Remains Unsolved. The boundary between productive self-improvement and model collapse is not well characterized. We lack theoretical guarantees for when verification-based self-training (where only correct solutions are retained) avoids collapse.

Research Directions. (1) Develop formal convergence proofs for verification-guided self-

improvement under distributional assumptions. (2) Design diversity-preserving self-training objectives that maintain coverage of the data distribution. (3) Establish information-theoretic bounds on the maximum self-improvement achievable without external supervision.

7.3 Multimodal Continual Learning

The Challenge. As LLMs evolve into multimodal foundation models processing text, images, audio, and video, continual learning must extend beyond language. The interaction between modalities introduces new forms of catastrophic interference—updating visual grounding can disrupt linguistic capabilities and vice versa.

Current Approaches. Adapter-based methods (?) offer modality-specific parameter isolation, where separate adapters handle different modalities. Mixture-of-experts architectures (?) provide natural modularity for multimodal routing. However, cross-modal representations in shared layers remain vulnerable to interference during sequential updates.

What Remains Unsolved. Cross-modal transfer during continual learning is poorly understood. When a model learns new visual concepts, how should associated linguistic knowledge update? The alignment between modality-specific and shared representations adds dimensionality to the stability-plasticity trade-off.

Research Directions. (1) Develop multimodal replay strategies that maintain cross-modal alignment. (2) Investigate whether modality-specific parameter isolation can preserve shared representations. (3) Create benchmarks for evaluating multimodal continual learning with controlled cross-modal dependencies.

7.4 Safe Continual Alignment

The Challenge. As models undergo continual learning, safety alignment—the product of careful RLHF (?) and constitutional AI (?) training—can degrade. This creates a tension between capability improvement and safety maintenance, often called the *alignment tax*: the cost of maintaining safety constraints during continual updates.

Current Approaches. Constitutional AI provides self-supervised alignment through critique-revision loops (?). Regularization toward safety-aligned parameter regions offers partial protection. Some work proposes “safety-locking” critical layers while allowing others to update freely.

What Remains Unsolved. No principled method guarantees that safety properties are preserved under arbitrary continual learning. ? demonstrate that even benign fine-tuning can degrade safety alignment, while ? show that RLHF itself reduces output diversity. The interaction between capability gain and alignment degradation is not formally characterized. Furthermore, as models self-improve, they may discover ways to circumvent alignment constraints that were not anticipated during initial training.

Research Directions. (1) Develop formal safety certificates that are maintained under bounded parameter updates. (2) Design alignment-aware continual learning objectives that jointly optimize for capability and safety. (3) Create adversarial evaluation suites that test alignment robustness under continual updates. (4) Investigate constitutional self-improvement where models continually refine their own alignment criteria.

7.5 Real-Time Learning at Deployment

The Challenge. The ultimate vision of continual learning requires models that update in real time while serving users. This poses severe infrastructure challenges: gradient computation interferes with inference throughput, weight updates must propagate across distributed serving systems, and consistency guarantees become non-trivial.

Current Approaches. Test-time training and in-context learning offer inference-time adapta-

tion without weight updates (?). Parameter-efficient methods like LoRA (?) reduce the computational burden of online updates. Some systems implement asynchronous training-serving pipelines where model checkpoints are periodically synchronized.

What Remains Unsolved. Latency constraints of real-time serving (typically <100ms) are fundamentally at odds with gradient computation. The staleness problem—serving predictions from outdated parameters while updates are being computed—lacks theoretical treatment. Quality assurance for continuous deployment (how to validate updates before they reach users) remains ad hoc.

Research Directions. (1) Develop efficient architectures that separate “fast” inference pathways from “slow” learning pathways. (2) Design consistency protocols for distributed model serving under continuous updates. (3) Create automated quality gates that validate continual updates against regression benchmarks before deployment.

7.6 Integration with Agent Frameworks

The Challenge. Autonomous agents built on LLMs require both continual learning (to accumulate knowledge from interactions) and self-improvement (to refine strategies based on outcomes) (?). This integration demands that CL and self-improvement operate not as independent modules but as tightly coupled components of a unified learning system.

Current Approaches. Reflexion (?) implements episodic self-improvement through verbal reflection, demonstrating that agents can learn from failures within a trajectory. Self-instruct (?) generates training data from model outputs, providing a self-improvement signal for downstream tasks. Long-running autonomous research frameworks that conduct literature review, experimentation, and writing exemplify the vision of agents that continually accumulate knowledge while iteratively improving their methods.

What Remains Unsolved. How should an agent decide *when* to consolidate episodic experience into parametric knowledge? The credit assignment problem—determining which past actions led to successful outcomes—is exacerbated in long-horizon agent tasks. Safety constraints must be maintained even as agents develop novel strategies through self-improvement.

Research Directions. (1) Design hierarchical memory architectures that support both short-term episodic and long-term parametric knowledge. (2) Develop meta-learning frameworks where agents improve their own learning algorithms through experience. (3) Create benchmarks for evaluating long-horizon agent self-improvement with safety constraints. (4) Investigate multi-agent continual learning where agents share and consolidate knowledge collectively.

7.7 Summary of Open Problems

The six challenges above share a common theme: the need for principled, theoretically grounded approaches that go beyond empirical heuristics. The most impactful near-term direction, in our assessment, is the integration of safe continual alignment with self-improvement—enabling models that grow more capable while provably maintaining safety guarantees.

8 Conclusion

This survey has provided the first unified treatment of continual learning and self-improvement in large language models, two research threads that have developed largely independently but are fundamentally convergent in their goals.

A Unified Taxonomy. We proposed a three-axis taxonomy organizing the field along *what* is learned (knowledge, skills, alignment), *how* learning occurs (external supervision, self-generated

signals, architectural adaptation), and *when* updates happen (offline, online, test-time). This framework reveals that methods traditionally categorized as “continual learning” or “self-improvement” often occupy adjacent positions in this space, differing primarily in their source of supervision rather than their underlying mechanisms.

Key Finding: Convergence of CL and Self-Improvement. Our analysis demonstrates that the most effective recent systems combine elements from both traditions. Self-improving models require continual learning mechanisms to avoid model collapse and catastrophic forgetting of previously acquired capabilities. Conversely, continual learning systems benefit from self-generated training signals to reduce dependence on external supervision. This convergence suggests that future progress will increasingly come from methods that jointly address both challenges.

State of the Art. The field has achieved notable successes: parameter-efficient methods enable continual adaptation with minimal forgetting; self-training with verification produces genuine capability improvement; and test-time adaptation offers immediate responsiveness without parameter modification. However, significant gaps remain. Theoretical understanding of self-improvement limits is nascent, safe continual alignment lacks formal guarantees, and multimodal continual learning is largely unexplored.

Call to Action. We identify three priorities for the research community. First, *unified benchmarks* that evaluate both retention and improvement simultaneously—current benchmarks measure one or the other, fragmenting evaluation. Second, *theoretical foundations* that characterize convergence conditions for self-improvement and scaling laws for continual learning capacity. Third, *safety-aware continual systems* that provably maintain alignment properties under ongoing learning, addressing the critical gap between capability advancement and safety assurance.

Looking Forward. The ultimate vision is of language models that perpetually learn and improve: absorbing new knowledge from their environment, refining their reasoning strategies through self-reflection, adapting to evolving user needs, and doing so while maintaining robust safety guarantees. Achieving this vision requires not just better algorithms but a fundamental rethinking of the training-deployment dichotomy. The model of the future is not trained and then deployed—it is deployed *in order to* learn. We hope this survey provides a foundation for researchers working toward that ambitious goal.

A Production Metrics

Both papers in the Deli AutoResearch series were autonomously generated by the **Deli AutoResearch** framework — an orchestration system that coordinates multiple AI agents to conduct literature research, write LaTeX content, verify citations, run experiments, and compile publications without human intervention during execution.

Table ?? provides a complete comparison of computational resources consumed across both papers.

Key Differences Between Paper #1 and #2. Paper #2 consumed approximately $4\times$ more tokens and $4\times$ more wall-clock time than Paper #1, primarily due to three additions not present in the first paper: (1) iterative peer review with three refinement cycles (vs. one), (2) multi-model experiments requiring $\sim 2,500$ API calls across 4 frontier models, and (3) vector figure generation via SVG $\rightarrow$ PDF conversion. Both papers achieved 100% citation verification; Paper #2 required

Table 8: Production statistics: Paper #1 (Autonomous Research Agents Survey) vs. Paper #2 (this paper, Continual Learning & Self-Improvement Survey). Totals reflect cumulative framework usage.

Metric	Paper #1	Paper #2	Total
<i>Compute & Time</i>			
Wall-clock time	2 h 22 min	~11 h	~13.4 h
Research iterations	5	10	15
Sub-agents spawned	12	18+	30+
Input tokens (est.)	320,000	1,100,000	1,420,000
Output tokens (est.)	140,000	480,000	620,000
Total tokens	460,000	1,580,000	2,040,000
Skill Hub skills invoked	1 (search_agent)	2 (search_agent, call_api)	2
API models used	1 (Opus 4.6)	5 (Opus, Sonnet, GPT-5.2, Gemini, DS)	5
<i>Citations & Verification</i>			
BibTeX entries	103	151	254
Citations web-verified	103 / 103 (100%)	151 / 151 (100%)	254 / 254 (100%)
Corrections applied	21	41	62
Hallucinated entries	0	0	0
<i>Paper Content</i>			
LaTeX source lines	2,070	2,100	4,170
PDF pages	42	47	89
Figures	0	8 (SVG → PDF vector)	8
Tables	3	9	12
Original experiments	0	1 (multi-model, 3 replications)	1
<i>Quality Assurance</i>			
Peer review score (initial)	6.0 / 10	6.0 / 10	—
Peer review score (final)	6.0 / 10	8.0 / 10	—
Review fix iterations	1	3	4
Pitfalls from Paper #1 applied	—	5 / 5	—

a second verification pass (using multi-model cross-checking) to resolve 41 corrections across 151 entries, compared to 21 corrections across 103 entries for Paper #1.

Iterative Process (Paper #2). This paper was produced in 9 iterations:

1. Template setup + Sections 1–2 + 54 initial citations
2. Core technical sections (Sections 3–4)
3. Application sections (Sections 5–6)
4. Final sections (7–8) + figure generation + bibliography expansion to 111 entries
5. Citation verification (20 sampled, 100% real, 2 venue corrections) + LLM peer review (6.0/10)
6. Review fixes + figure vectorization + appendix + initial deployment
7. **V2:** Full citation verification (111 verified, 10 corrected) + 39 new references + quantitative tables + formal propositions + CL×SI interaction section + re-review (7.5/10)
8. **V2:** Additional figures (interaction matrix, performance profiles) + impossibility theorem + taxonomy gap analysis + deployment
9. **V3:** Multi-model experiment (4 models × 3 conditions × 3 replications) + replication study + taxonomy limitations + model merging coverage + prompt-replay theory + claim calibration + re-review (8.0/10) + deployment

10. **V4:** Full 151/151 citation verification (multi-model cross-check) + 41 corrections applied + arXiv compliance (AI tools moved to footnote) + all undefined references fixed + final deployment

Lessons Learned (Cumulative). Key operational insights accumulated across both papers:

- Compiling LaTeX in Iteration 1 catches environment issues early (missing .sty files, font problems)
- Incremental citation verification prevents error accumulation; batch verification is more efficient but riskier
- Vector figures (SVG $\rightarrow$ PDF via `cairosvg`) are essential for print quality; raster fallbacks should be avoided
- LLM peer review at 50% completion enables earlier course correction than post-hoc review
- Production metrics should be tracked from the start, not reconstructed post-hoc
- Original experiments—even small pilot studies—significantly differentiate a survey from a literature review (reviewer feedback: +1 point)
- Softening formal claims (“conjecture” vs. “theorem”) is preferred when full proofs are not provided
- Multi-trial replication with variance estimates is more convincing than larger single-trial n for pilot studies

References

- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, et al. GPT-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Hongjoon Ahn, Sungmin Cha, Donggyu Lee, and Taesup Moon. Uncertainty-based continual learning with adaptive regularization. In *Advances in Neural Information Processing Systems*, 2019.
- Sina Alemohammad, Josue Casco-Rodriguez, Lorenzo Luber, Ahmed Babaei, Daniel LeJeune, Ali Siahkoobi, and Richard G Baraniuk. Self-consuming generative models go MAD. In *International Conference on Learning Representations*, 2024.
- Anthropic. The Claude 3 model family: A new standard for intelligence. *Anthropic Technical Report*, 2024.
- Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, et al. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*, 2022a.
- Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, et al. Constitutional AI: Harmlessness from AI feedback. *arXiv preprint arXiv:2212.08073*, 2022b.
- Dan Biderman, Jacob Portes, Jose Javier Gonzalez Ortiz, Philip Greengard, Allen Gersho, Jonathan Frankle, Danqi Xiao, and Alessandro Stolfo. Lora learns less and forgets less. *arXiv preprint arXiv:2405.09673*, 2024.

- Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George Bm van den Driessche, Jean-Baptiste Lespiau, Bogdan Damoc, Aidan Clark, et al. Improving language models by retrieving from trillions of tokens. In *International Conference on Machine Learning*, 2022.
- Bradley Brown, Jordan Juravsky, Ryan Ehrlich, Ronald Clark, Quoc V Le, Christopher Ré, and Azalia Mirhoseini. Large language monkeys: Scaling inference compute with repeated sampling. *arXiv preprint arXiv:2407.21787*, 2024.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. In *Advances in Neural Information Processing Systems*, 2020.
- Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Weak-to-strong generalization: Eliciting strong capabilities with weak supervision. *arXiv preprint arXiv:2312.09390*, 2023.
- Pietro Buzzega, Matteo Boschini, Angelo Porrello, Davide Abati, and Simone Calderara. Dark experience for general continual learning: a strong, simple baseline. In *Advances in Neural Information Processing Systems*, 2020.
- Zixiang Chen, Yihe Deng, Huizhuo Yuan, Kaixuan Ji, and Quanquan Gu. Self-play fine-tuning converts weak language models to strong language models. In *International Conference on Machine Learning*, 2024.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Matthias De Lange, Rahaf Aljundi, Marc Masana, Sarah Parisot, Xu Jez, Ales Leonardis, Gregory Slabaugh, and Tinne Tuytelaars. A continual learning survey: Defying forgetting in classification tasks. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 44(7):3366–3385, 2021.
- DeepSeek-AI. DeepSeek-V2: A strong, economical, and efficient mixture-of-experts language model. *arXiv preprint arXiv:2405.04434*, 2024.
- DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- Elvis Dohmatob, Yunzhen Feng, Arjun Subramonian, and Julia Kempe. A tale of tails: Model collapse as a change of scaling laws. *arXiv preprint arXiv:2402.07043*, 2024.
- Hanze Dong, Wei Xiong, Bo Pang, Haoxiang Wang, Han Zhao, Yingbo Zhou, Nan Jiang, Doyen Sahoo, Caiming Xiong, and Tong Zhang. RLHF workflow: From reward modeling to online RLHF. *arXiv preprint arXiv:2405.07863*, 2024.
- Yilun Du, Shuang Li, Antonio Torralba, Joshua B Tenenbaum, and Igor Mordatch. Improving factuality and reasoning in language models through multiagent debate. *arXiv preprint arXiv:2305.14325*, 2024.
- Logan Engstrom, Andrew Ilyas, Shibani Santurkar, Dimitris Tsipras, Fangzhou Wen, and Aleksander Madry. Data selection for language models via importance resampling. In *Advances in Neural Information Processing Systems*, 2024.

- William Fedus, Barret Zoph, and Noam Shazeer. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *Journal of Machine Learning Research*, 23(120):1–39, 2022.
- Chelsea Finn, Pieter Abbeel, and Sergey Levine. Model-agnostic meta-learning for fast adaptation of deep networks. In *International Conference on Machine Learning*, 2017.
- Matthias Gerstgrasser, Sam Schoenholz, Ari Stern, Michael Deweese, and Joel Dyer. Is model collapse inevitable? breaking the curse of recursion by accumulating real and synthetic data. *arXiv preprint arXiv:2404.01413*, 2024.
- Caglar Gulcehre, Tom Le Paine, Srivatsan Srinivasan, Dmitrii Kochkov, Lasse Ke, Konrad Zolna, et al. Reinforced self-training (ReST) for language modeling. In *International Conference on Machine Learning*, 2023.
- Suriya Gunasekar, Yi Zhang, Jyoti Anber, Girish Menber, Adele Del Giorno, Nikhil Gopi, Mojan Javadi, Saurabh Kauffman, Gustavo de Rosa, Olli Saarikivi, et al. Textbooks are all you need. *arXiv preprint arXiv:2306.11644*, 2023.
- Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Mingwei Chang. REALM: Retrieval-augmented language model pre-training. In *International Conference on Machine Learning*, 2020.
- Danny Hernandez, Jared Kaplan, Tom Henighan, and Sam McCandlish. Scaling laws for transfer. *arXiv preprint arXiv:2102.01293*, 2021.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.
- Arian Hosseini, Xingdi Yuan, Nikolay Malber, Aaron Courville, and Alessandro Sordoni. V-STaR: Training verifiers for self-taught reasoners. *arXiv preprint arXiv:2402.06457*, 2024.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.
- Tongtian Huang, Delin Chu, Ziqian Zhuge, et al. Continual learning for large language models: A survey. *arXiv preprint arXiv:2402.01364*, 2024.
- Gabriel Ilharco, Marco Tulio Ribeiro, Mitchell Wortsman, Suchin Gururangan, Ludwig Schmidt, Hannaneh Hajishirzi, and Ali Farhadi. Editing models with task arithmetic. In *International Conference on Learning Representations*, 2023.
- Joel Jang, Seonghyeon Ye, Sohee Yang, Joongbo Shin, Janghoon Han, Gyeonghun Kim, Stanley Jungkyu Choi, and Minjoon Seo. Towards continual knowledge learning of language models. In *International Conference on Learning Representations*, 2022.
- Xisen Jin, Dejiao Zhang, Henghui Zhu, Wei Xiao, Shang-Wen Li, Xiaokai Wei, Andrew Arnold, and Xiang Ren. Lifelong pretraining: Continually adapting language models to emerging corpora. In *North American Chapter of the Association for Computational Linguistics*, 2022.
- Zixuan Ke, Yijia Shao, Haowei Lin, Tatsuya Konishi, Gyuhak Kim, and Bing Liu. Continual pre-training of language models. In *International Conference on Learning Representations*, 2023.

- Akbir Khan, John Hughes, Dan Roth, Kanishka Misra, Ethan Perez, and Samuel R Bowman. Debating with more persuasive LLMs leads to more truthful answers. *arXiv preprint arXiv:2402.06782*, 2024.
- Robert Kirk, Ishita Mediratta, Christoforos Nalmpantis, Jelena Luketina, Eric Hambro, Edward Grefenstette, and Roberta Raileanu. Understanding the effects of RLHF on LLM generalisation and diversity. *arXiv preprint arXiv:2310.06452*, 2024.
- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, et al. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 2017.
- Zhizhong Li and Derek Hoiem. Learning without forgetting. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 40(12):2935–2947, 2017.
- Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *International Conference on Learning Representations*, 2024.
- Bill Yuchen Lin, Abhilasha Ravichander, Ximing Lu, Nouha Dziri, Melanie Sclar, Khyathi Chandu, Chandra Bhagavatula, and Yejin Choi. The unlocking spell on base LLMs: Rethinking alignment via in-context learning. *arXiv preprint arXiv:2312.01552*, 2024.
- David Lopez-Paz and Marc’Aurelio Ranzato. Gradient episodic memory for continual learning. In *Advances in Neural Information Processing Systems*, 2017.
- Haipeng Luo, Qingfeng Sun, Can Xu, Pu Zhao, Jianguang Lou, Chongyang Tao, Xiubo Geng, Qingwei Lin, Shifeng Chen, and Dongmei Zhang. WizardMath: Empowering mathematical reasoning for large language models via reinforced evol-instruct. In *International Conference on Learning Representations*, 2024a.
- Yun Luo, Zhen Yang, Fandong Meng, Yafu Li, Jie Zhou, and Yue Zhang. An empirical study of catastrophic forgetting in large language models during continual fine-tuning. *arXiv preprint arXiv:2308.08747*, 2024b.
- Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, et al. Self-refine: Iterative refinement with self-feedback. In *Advances in Neural Information Processing Systems*, 2023.
- Arun Mallya and Svetlana Lazebnik. Packnet: Adding multiple tasks to a single network by iterative pruning. In *IEEE Conference on Computer Vision and Pattern Recognition*, 2018.
- Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems*, 2022.
- Kevin Meng, Arnab Sen Sharma, Alex Andonian, Yonatan Belinkov, and David Bau. Mass-editing memory in a transformer. In *International Conference on Learning Representations*, 2023.
- Eric Mitchell, Charles Lin, Antoine Bosselut, Chelsea Finn, and Christopher D Manning. Fast model editing at scale. In *International Conference on Learning Representations*, 2022.
- OpenAI. GPT-4o system card. *OpenAI Technical Report*, 2024a.

- OpenAI. Learning to reason with LLMs. *OpenAI Blog*, 2024b.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems*, 2022.
- Richard Yuanzhe Pang, Weizhe Yuan, Kyunghyun Cho, He He, Sainbayar Sukhbaatar, and Jason Weston. Iterative reasoning preference optimization. *arXiv preprint arXiv:2404.19733*, 2024.
- German I Parisi, Ronald Kemker, Jose L Part, Christopher Kanan, and Stefan Wermter. Continual lifelong learning with neural networks: A review. *Neural Networks*, 113:54–71, 2019.
- Jonas Pfeiffer, Aishwarya Kamath, Andreas Rücklé, Kyunghyun Cho, and Iryna Gurevych. AdapterFusion: Non-destructive task composition for transfer learning. In *Conference of the European Chapter of the Association for Computational Linguistics*, 2021.
- Xiangyu Qi, Yi Zeng, Tinghao Xie, Pin-Yu Chen, Ruoxi Jia, Prateek Mittal, and Peter Henderson. Fine-tuning aligned language models compromises safety, even when users do not intend to! *arXiv preprint arXiv:2310.03693*, 2024.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D Manning, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems*, 2023.
- Anastasia Razdaibiedina, Yuning Mao, Rui Hou, Madian Khabisa, Mike Lewis, and Amjad Almahairi. Progressive prompts: Continual learning for language models. In *International Conference on Learning Representations*, 2023.
- Andrei A Rusu, Neil C Rabinowitz, Guillaume Desjardins, Hubert Sober, Koray Kavukcuoglu, and Razvan Pascanu. Progressive neural networks. *arXiv preprint arXiv:1606.04671*, 2016.
- Thomas Scialom, Tuhin Chakrabarty, and Smaranda Muresan. Fine-tuned language models are continual learners. In *Conference on Empirical Methods in Natural Language Processing*, 2022.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Mingchuan Zhang, Yong Li, Yu Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- Noam Shazeer, Azalia Miber, Krzysztof Parameswar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. In *International Conference on Learning Representations*, 2017.
- Hanul Shin, Jung Kwon Lee, Jaehong Kim, and Jiwon Kim. Continual learning with deep generative replay. In *Advances in Neural Information Processing Systems*, 2017.
- Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems*, 2023.
- Ilya Shumailov, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. Ai models collapse when trained on recursively generated data. *Nature*, 631:755–759, 2024.

- David Silver, Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, et al. Mastering the game of Go without human knowledge. *Nature*, 550:354–359, 2017.
- Avi Singh, John D Co-Reyes, Rishabh Agarwal, Ankesh Anand, Piyush Patil, Xavier Garcia, Peter J Liu, James Harrison, Jaehoon Lee, Kelvin Xu, et al. Beyond human data: Scaling self-training for problem-solving with language models. *Transactions on Machine Learning Research*, 2024.
- James Seale Smith, Leonid Karlinsky, Vyshnavi Gutta, Paola Cascante-Bonilla, Donghyun Kim, Assaf Arbelle, Rameswar Panda, Rogerio Feris, and Zsolt Kira. CODA-Prompt: Continual decomposed attention-based prompting for rehearsal-free continual learning. In *IEEE Conference on Computer Vision and Pattern Recognition*, 2023.
- Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. Scaling LLM test-time compute optimally can be more effective than scaling model parameters. In *International Conference on Learning Representations*, 2025.
- Nisan Stiennon, Long Ouyang, Jeffrey Wu, Daniel Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford, Dario Amodei, and Paul F Christiano. Learning to summarize with human feedback. In *Advances in Neural Information Processing Systems*, 2020.
- Sebastian Thrun. *Lifelong Learning Algorithms*. Springer, 1998.
- Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. LLaMA: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- Gido M van de Ven and Andreas S Tolias. Three scenarios for continual learning. In *NeurIPS Continual Learning Workshop*, 2019.
- Tongtong Wang, Xinhao Li, Qi Zhu, Jingbo Lian, et al. A comprehensive survey of continual learning for large language models. *arXiv preprint arXiv:2404.16789*, 2024a.
- Xiao Wang, Yuansen Chen, Naiheng Zhu, Jielong Li, Qiang Gao, and Liang Hou. TRACE: A comprehensive benchmark for continual learning in large language models. *arXiv preprint arXiv:2310.06762*, 2024b.
- Yifan Wang, Jinguo Ye, Wanjuan Wang, et al. InsCL: A data-efficient continual learning paradigm for fine-tuning large language models with instructions. *arXiv preprint arXiv:2403.11435*, 2024c.
- Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, Alisa Liu, Noah A Smith, Daniel Khashabi, and Hannaneh Hajishirzi. Self-instruct: Aligning language models with self-generated instructions. In *Association for Computational Linguistics*, 2023.
- Zifeng Wang, Zizhao Zhang, Sayna Ebrahimi, Ruoxi Sun, Han Zhang, Chen-Yu Lee, Xiaoqi Ren, Guolong Su, Vincent Perot, Jennifer Dy, and Tomas Pfister. DualPrompt: Complementary prompting for rehearsal-free continual learning. In *European Conference on Computer Vision*, 2022a.
- Zifeng Wang, Zizhao Zhang, Chen-Yu Lee, Han Zhang, Ruoxi Sun, Xiaoqi Ren, Guolong Su, Vincent Perot, Jennifer Dy, and Tomas Pfister. Learning to prompt for continual learning. In *IEEE Conference on Computer Vision and Pattern Recognition*, 2022b.

- Yue Wu, Zhiqing Sun, Huizhuo Yuan, Kaixuan Ji, Yiming Yang, and Quanquan Gu. Self-play preference optimization for language model alignment. *arXiv preprint arXiv:2405.00675*, 2024.
- Sang Michael Xie, Shibani Santurkar, Tengyu Ma, and Percy Liang. DoReMi: Optimizing data mixtures speeds up language model pretraining. In *Advances in Neural Information Processing Systems*, 2023.
- Wei Xiong, Hanze Dong, Chenlu Ye, Ziqi Wang, Han Zhong, Heng Ji, Nan Jiang, and Tong Zhang. Iterative preference learning from human feedback: Bridging theory and practice for RLHF under KL-constraint. *arXiv preprint arXiv:2312.11456*, 2024.
- Can Xu, Qingfeng Sun, Kai Zheng, Xiubo Geng, Pu Zhao, Jiazhan Feng, Chongyang Tao, and Daxin Jiang. WizardLM: Empowering large language models to follow complex instructions. In *International Conference on Learning Representations*, 2024.
- Prateek Yadav, Derek Tam, Leshem Choshen, Colin Raffel, and Mohit Bansal. TIES-Merging: Resolving interference when merging models. In *Advances in Neural Information Processing Systems*, 2023.
- Prateek Yadav, Derek Tam, Leshem Choshen, Colin Raffel, and Mohit Bansal. Ties-merging: Resolving interference when merging models. In *NeurIPS*, 2024.
- Le Yu, Bowen Yu, Haiyang Yu, Fei Huang, and Yongbin Li. Language models are Super Mario: Absorbing abilities from homologous models as a free lunch. In *International Conference on Machine Learning*, 2024.
- Weizhe Yuan, Richard Yuanzhe Pang, Kyunghyun Cho, Sainbayar Sukhbaatar, Jing Xu, and Jason Weston. Self-rewarding language models. *arXiv preprint arXiv:2401.10020*, 2024.
- Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah D Goodman. STaR: Bootstrapping reasoning with reasoning. In *Advances in Neural Information Processing Systems*, 2022.
- Friedemann Zenke, Ben Poole, and Surya Ganguli. Continual learning through synaptic intelligence. In *International Conference on Machine Learning*, 2017.
- Ziheng Zhang, Luyang Lin, Dongyuan Cao, Wanjuan Wang, and Jingwei Liao. Continual instruction tuning for large language models. *arXiv preprint arXiv:2403.12249*, 2024.
- Daniel M Ziegler, Nisan Stiennon, Jeffrey Wu, Tom B Brown, Alec Radford, Dario Amodei, Paul Christiano, and Geoffrey Irving. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593*, 2019.